

SoK: Computer-Aided Cryptography

Manuel Barbosa*, Gilles Barthe^{†‡}, Karthik Bhargavan[§], Bruno Blanchet[§], Cas Cremers[¶], Kevin Liao^{†||}, Bryan Parno**

*University of Porto (FCUP) and INESC TEC, [†]Max Planck Institute for Security & Privacy, [‡]IMDEA Software Institute,

[§]INRIA Paris, [¶]CISPA Helmholtz Center for Information Security, ^{||}MIT, **Carnegie Mellon University

Abstract—Computer-aided cryptography is an active area of research that develops and applies formal, machine-checkable approaches to the design, analysis, and implementation of cryptography. We present a cross-cutting systematization of the computer-aided cryptography literature, focusing on three main areas: (i) design-level security (both symbolic security and computational security), (ii) functional correctness and efficiency, and (iii) implementation-level security (with a focus on digital side-channel resistance). In each area, we first clarify the role of computer-aided cryptography—how it can help and what the caveats are—in addressing current challenges. We next present a taxonomy of state-of-the-art tools, comparing their accuracy, scope, trustworthiness, and usability. Then, we highlight their main achievements, trade-offs, and research challenges. After covering the three main areas, we present two case studies. First, we study efforts in combining tools focused on different areas to consolidate the guarantees they can provide. Second, we distill the lessons learned from the computer-aided cryptography community’s involvement in the TLS 1.3 standardization effort. Finally, we conclude with recommendations to paper authors, tool developers, and standardization bodies moving forward.

I. INTRODUCTION

Designing, implementing, and deploying cryptographic mechanisms is notoriously hard to get right, with high-profile design flaws, devastating implementation bugs, and side-channel vulnerabilities being regularly found even in widely deployed mechanisms. Each step is highly involved and fraught with pitfalls. At the design level, cryptographic mechanisms must achieve specific security goals against some well-defined class of attackers. Typically, this requires composing a series of sophisticated building blocks—abstract constructions make up primitives, primitives make up protocols, and protocols make up systems. At the implementation level, high-level designs are then fleshed out with concrete functional details, such as data formats, session state, and programming interfaces. Moreover, implementations must be optimized for interoperability and performance. At the deployment level, implementations must also account for low-level threats that are absent at the design level, such as side-channel attacks.

Attackers are thus presented with a vast attack surface: They can break high-level designs, exploit implementation bugs, recover secret material via side-channels, or any combination of the above. Preventing such varied attacks on complex cryptographic mechanisms is a challenging task, and existing methods are hard-pressed to do so. Pen-and-paper security proofs often consider pared-down “cores” of cryptographic mechanisms to simplify analysis, yet remain highly complex and error-prone; demands for aggressively optimized implementations greatly increase the risks of introducing bugs,

which are difficult to catch by code testing or auditing; ad-hoc constant-time coding recipes for mitigating side-channel attacks are tricky to implement, and yet may not cover the whole gamut of leakage channels exposed in deployment. Unfortunately, the current modus operandi—relying on a select few cryptography experts armed with rudimentary tooling to vouch for security and correctness—simply cannot keep pace with the rate of innovation and development in the field.

Computer-aided cryptography, or CAC for short, is an active area of research that aims to address these challenges. It encompasses formal, machine-checkable approaches to designing, analyzing, and implementing cryptography; the variety of tools available address different parts of the problem space. At the design level, tools can help manage the complexity of security proofs, even revealing subtle flaws or as-yet-unknown attacks in the process. At the implementation level, tools can guarantee that highly optimized implementations behave according to their design specifications on all possible inputs. At the deployment level, tools can check that implementations correctly protect against classes of side-channel attacks. Although individual tools may only address part of the problem, when combined, they can provide a high degree of assurance.

Computer-aided cryptography has already fulfilled some of these promises in focused but impactful settings. For instance, computer-aided security analyses were influential in the recent standardization of TLS 1.3 [1]–[4]. Formally verified code is also being deployed at Internet-scale—components of the HACL* library [5] are being integrated into Mozilla Firefox’s NSS security engine, elliptic curve code generated using the Fiat Cryptography framework [6] has populated Google’s BoringSSL library, and EverCrypt [7] routines are used in the Zinc crypto library for the Linux kernel. In light of these successes, there is growing enthusiasm for computer-aided cryptography. This is reflected in the rapid emergence of a dynamic community comprised of theoretical and applied cryptographers, cryptography engineers, and formal methods practitioners. Together, the community aims to achieve broader adoption of computer-aided cryptography, blending ideas from many fields, and more generally, to contribute to the future development of cryptography.

At the same time, computer-aided cryptography risks becoming a victim of its own success. Trust in the field can be undermined by difficulties in understanding the guarantees and fine-print caveats of computer-aided cryptography artifacts. The field is also increasingly broad, complex, and rapidly evolving, so no one has a complete understanding of every facet. This can make it difficult for the field to develop and

address pressing challenges, such as the expected transition to post-quantum cryptography and scaling from lower-level primitives and protocols to whole cryptographic systems.

Given these concerns, the purpose of this SoK is three-fold:

- 1) We clarify the current capabilities and limitations of computer-aided cryptography.
- 2) We present a taxonomy of computer-aided cryptography tools, highlighting their main achievements and important trade-offs between them.
- 3) We outline promising new directions for computer-aided cryptography and related areas.

We hope this will help non-experts better understand the field, point experts to opportunities for improvement, and showcase to stakeholders (e.g., standardization bodies and open source projects) the many benefits of computer-aided cryptography.

A. Structure of the Paper

The subsequent three sections expand on the role of computer-aided cryptography in three main areas: Section II covers how to establish design-level security guarantees, using both symbolic and computational approaches; Section III covers how to develop functionally correct and efficient implementations; Section IV covers how to establish implementation-level security guarantees, with a particular focus on protecting against digital side-channel attacks.

We begin each section with a critical review of the area, explaining why the considered guarantees are important, how current tools and techniques outside CAC may fail to meet these guarantees, how CAC can help, the fine-print caveats of using CAC, and necessary technical background. We then taxonomize state-of-the-art tools based on criteria along four main categories: *accuracy* (*A*), *scope* (*S*), *trust* (*T*), and *usability* (*U*). For each criterion, we label them with one or more categories, explain their importance, and provide some light discussion about tool support for them. The ensuing discussion highlights broader points, such as main achievements, important takeaways, and research challenges. Finally, we end each section with references for further reading. Given the amount of material we cover, we are unable to be exhaustive in each area, but we still point to other relevant lines of work.

Sections V and VI describe two case studies. Our first case study (Section V) examines how to combine tools that address different parts of the problem space and consolidate their guarantees. Our second case study (Section VI) distills the lessons learned from the computer-aided cryptography community's involvement in the TLS 1.3 standardization effort.*

Finally, in Section VII, we offer recommendations to paper authors, tool developers, and standardization bodies on how to best move the field of computer-aided cryptography forward.

II. DESIGN-LEVEL SECURITY

In this section, we focus on the role of computer-aided cryptography in establishing design-level security guarantees. Over the years, two flavors of design-level security have been developed in two largely separate communities: symbolic security (in the formal methods community) and computational

security (in the cryptography community). This has led to two complementary strands of work, so we cover them both.

A. Critical Review

Why is design-level security important? Validating cryptographic designs through mathematical arguments is perhaps the only way to convincingly demonstrate their security against entire classes of attacks. This has become standard practice in cryptography, and security proofs are necessary for any new standard. This holds true at all levels: primitives, protocols, and systems. When using a lower-level component in a larger system, it is crucial to understand what security notion and adversarial model the proof is relative to. Similar considerations apply when evaluating the security of a cryptographic system relative to its intended deployment environment.

How can design-level security fail? The current modus operandi for validating the security of cryptographic designs using pen-and-paper arguments is alarmingly fragile. This is for two main reasons:

- **Erroneous arguments.** Writing security arguments is tedious and error-prone, even for experts. Because they are primarily done on pen-and-paper, errors are difficult to catch and can go unnoticed for years.
- **Inappropriate modeling.** Even when security arguments are correct, attacks can lie outside the model in which they are established. This is a known and common pitfall: To make (pen-and-paper) security analysis tractable, models are often heavily simplified into a cryptographic core that elides many details about cryptographic designs and attacker capabilities. Unfortunately, attacks are often found outside of this core.

How are these failures being addressed outside CAC?

To minimize erroneous arguments, cryptographers have devised a number of methodological frameworks for security analysis (e.g., the code-based game playing [8] and universal composability [9] frameworks). The high-level goal of these frameworks is to decompose security arguments into simpler arguments that are easier to get right and then smoothly combine the results. Still, pen-and-paper proofs based on these methodologies remain complex and error-prone, which has led to suggestions of using computer-aided tools [10]. → Halevi 05.

To reduce the risks of inappropriate modeling, real-world provable security [11]–[13]* advocates making security arguments in more accurate models of cryptographic designs and adversarial capabilities. Unfortunately, the added realism comes with greater complexity, complicating security analysis.

How can computer-aided cryptography help? Computer-aided cryptography tools are effective for detecting flaws in cryptographic designs and for managing the complexity of security proofs. They crystallize the benefits of on-paper methodologies and of real-world provable security. They also deliver trustworthy analyses for complex designs that are beyond reach of pen-and-paper analysis.

What are the fine-print caveats? Computer-aided security proofs are only as good as the statements being proven. However, understanding these statements can be challenging. Most security proofs rely on implicit assumptions; without

* Read more

* Added the paper at the end of this one

proper guidance, reconstructing top-level statements can be challenging, even for experts. (As an analogy, it is hard even for a talented mathematician to track all dependencies in a textbook.) Finally, as with any software, tools may have bugs.

What background do I need to know about symbolic security? The symbolic model is an abstract model for representing and analyzing cryptographic protocols. Messages (e.g., keys, nonces) are represented symbolically as *terms* (in the parlance of formal logic). Typically, terms are atomic data, meaning that they cannot be split into, say, component bitstrings. Cryptographic primitives are modeled as black-box functions over terms related by a set of mathematical identities called an *equational theory*. For example, symmetric encryption can be modeled by the black-box functions Enc and Dec related by the following equational theory: $\text{Dec}(\text{Enc}(m, k), k) = m$. This says that decrypting the ciphertext $\text{Enc}(m, k)$ using the key k recovers the original plaintext m .

An adversary is restricted to compute (i.e., derive new terms contributing to its *knowledge set*) using only the specified primitives and equational theory. Equational theories are thus important for broadening the scope of analysis—ignoring valid equations implicitly weakens the class of adversaries considered. In the example above, m and k are atomic terms, and so equipped with only the given identity, an adversary can decrypt a ciphertext only if it has knowledge of the *entire* secret key. Such simplifications enable modeling and verifying protocols using symbolic logic. Symbolic tools are thus well-suited to automatically searching for and unveiling logical flaws in complex cryptographic protocols and systems.

Symbolic security properties come in two main flavors: trace properties and equivalence properties. Trace properties state that a bad event never occurs on any execution trace. For example, a protocol preserves trace-based secrecy if, for any execution trace, secret data is not in the adversarial knowledge set. On the other hand, equivalence properties state that an adversary is unable to distinguish between two protocols, often with one being the security specification. Equivalence properties typically cannot be (naturally or precisely) expressed as trace properties. For example, a protocol preserves indistinguishability-based secrecy if the adversary cannot differentiate between a trace with the real secret and a trace with the real secret replaced by a random value.

What background do I need to know about computational security? In the computational model, messages are bitstrings, cryptographic primitives are probabilistic algorithms on bitstrings, and adversaries are probabilistic Turing machines. For example, symmetric encryption can be modeled by a triple of algorithms (Gen, Enc, Dec). The probabilistic key generation algorithm Gen outputs a bitstring k . The encryption (decryption) algorithm Enc (Dec) takes as input a key k and a plaintext m (ciphertext c), and outputs a ciphertext c (plaintext m). The basic correctness property that must hold for every key k output by Gen and every message m in the message space is $\text{Dec}(\text{Enc}(m, k), k) = m$. Because keys are bitstrings in this model, knowing bits of an encryption key reduces the computational resources required to decrypt a ciphertext.

Computational security properties are also probabilistic and can be characterized along two axes: game-based or simulation-based, and concrete or asymptotic.

Game-based properties specify a probabilistic experiment called a “game” between a challenger and an adversary, and an explicit goal condition that the adversary must achieve to break a scheme. Informally, security statements say: For all adversaries, the probability of achieving the goal condition does not exceed some threshold. The specific details, e.g., the adversary’s computational resources and the threshold, depend on the choice of concrete or asymptotic security.

A core proof methodology for game-based security is *game hopping*. In the originally specified game, the adversary’s success probability may be unknown. Thus, we proceed by step-wise transforming the game until reaching one in which the success probability can be computed. We also bound the increases in the success probability from the game transformations, often by reducing to an assumed hard problem (e.g., the discrete log or RSA problems). We can then deduce a bound on the adversary’s success probability in the original game. The interested reader can see the tutorials on game hopping by Shoup [14] and Bellare and Rogaway [8].

Simulation-based properties specify two probabilistic experiments: The “real” game runs the scheme under analysis. The “ideal” game runs an idealized scheme that does not involve any cryptography, but instead runs a trusted third-party called an *ideal functionality*, which serves as the security specification. Informally, security statements say: For all adversaries in the real game, there exists a *simulator* in the ideal game that can translate any attack on the real scheme into an attack on the ideal functionality. Because the ideal functionality is secure by definition, the real scheme must also be secure. In general, simulation-based proofs tend to be more complicated than game-based proofs, but importantly they support composition theorems that allow analyzing complex constructions in a modular way from simpler building blocks. The interested reader can see the tutorial on simulation-based proofs by Lindell [15].

Concrete security quantifies the security of a scheme by bounding the maximum success probability of an adversary given some upper bound on running time. A scheme is (t, ϵ) -secure if every adversary running for time at most t succeeds in breaking the scheme with probability at most ϵ . In contrast, asymptotic security views the running time of the adversary and its success probability as functions of some security parameter (e.g., key length), rather than as concrete numbers. A scheme is secure if every probabilistic polynomial time adversary succeeds in breaking the scheme with negligible probability (i.e., with probability asymptotically less than all inverse polynomials in the security parameter).

Of these different security properties, we note that computer-aided security proofs have primarily focused on game-based, concrete security. Work on mechanizing simulation-based proofs is relatively nascent; asymptotic security is the prevailing paradigm in cryptography, but by proving concrete security, asymptotic security follows *a fortiori*.

Tool	Unbound	Trace	Equiv	Eq-thy	State	Link
CPSA [▷] [16]	●	●	○	○	●	○
F7 [○] [17]	●	●	○	●	●	●
⊢FS [○] [18]	●	●	○	○	●	●
Maude-NPA [▷] [19]	●	●	● ^d	○	○	○
ProVerif ^{*†} [20]	●	●	● ^d	○	○	○
⊢fs2pv ^{○†} [21]	●	●	○	○	○	●
⊢GSVerif ^{*†} [22]	●	●	○	○	○	○
⊢ProVerif-ATP ^{*†} [23]	●	●	○	●	○	○
⊢StatVerif ^{*†} [24]	●	●	● ^d	○	●	○
Scyther [▷] [25]	●	●	○	○	○	○
sclyther-proof ^{▷‡§} [26]	●	●	○	○	○	○
Tamarin ^{*‡} [27]	●	●	● ^d	●	●	○
⊢SAPIC [*] [28]	●	●	○	●	●	○
CF-AtSe [▷] [29]	○	●	○	●	●	○
OFMC ^{▷†} [30]	○	●	○	○	●	○
SATMC [▷] [31]	○	●	○	○	●	○
AKISS ^{*†} [32]	○	○	● ^t	●	●	○
APTE [*] [33]	○	○	● ^t	○	●	○
DEEPSEC [*] [34]	○	○	● ^t	○	●	○
SAT-Equiv [*] [35]	○	○	● ^t	○	○	○
SPEC ^{*‡§} [36]	○	○	○ ^o	○	○	○

Specification language	Miscellaneous symbols
▷ – security protocol notation	⊢ – previous tool extension
* – process calculus	† – abstractions
* – multiset rewriting	‡ – interactive mode
◊ – general programming language	§ – independent verifiability

Equational theories (Eq-thy)	Equivalence properties (Equiv)
● – with AC axioms	t – trace equivalence
○ – without AC axioms	o – open bisimilarity
○ – fixed	d – diff equivalence

TABLE I

OVERVIEW OF TOOLS FOR SYMBOLIC SECURITY ANALYSIS. SEE SECTION II-B FOR MORE DETAILS ON COMPARISON CRITERIA.

B. Symbolic Tools: State of the Art

Table I presents a taxonomy of modern, general-purpose symbolic tools. Tools are listed in three groups (demarcated by dashed lines): unbounded trace-based tools, bounded trace-based tools, and equivalence-based tools; within each group, top-level tools are listed alphabetically. Tools are categorized as follows, annotated with the relevant criteria (A, S, T, U) described in the introduction. Note that the capabilities of symbolic tools are more nuanced than what is reflected in the table—the set of examples that tools can handle varies even if they support the same features according to the table.

Unbounded number of sessions (A). Can the tool analyze an unbounded number of protocol sessions? There exist protocols that are secure when at most N sessions are considered, but become insecure with more than N sessions [37]. Bounded tools (○) explicitly limit the analyzed number of sessions and do not consider attacks beyond the cut-off. Unbounded tools (●) can prove the absence of attacks within the model, but at the cost of undecidability [38].

In practice, modern unbounded tools typically substantially outperform bounded tools even for a small number of sessions, and therefore enable the analysis of more complex models. This is because bounded tools are a bit naive in their exploration of the state space, basically enumerating options (but exploiting some symmetry). They therefore typically grow exponentially in the number of sessions. The unbounded tools inherently need to be “more clever” to even achieve unbounded analysis. While their algorithms are more complex, when

they work (i.e., terminate), the analysis is independent of the number of sessions.

Trace properties (S). Does the tool support verification of trace properties?

Equivalence properties (S). Does the tool support verification of equivalence properties? There are several different equivalence notions used in current tools. Here, we provide some high-level intuition, but for a more formal treatment, see the survey by Delaune and Hirschi [39].

Trace equivalence (t) means that, for each trace of one protocol, there exists a corresponding trace of the other protocol, such that the messages exchanged in these two traces are indistinguishable. This is the weakest equivalence notion, roughly meaning that it can express the most security properties. (The other stronger notions are often intermediate steps towards proving trace equivalence.) It is also arguably the most natural for formalizing privacy properties.

Open bisimilarity (o) is a strictly stronger notion that captures the knowledge of the adversary by pairs of symbolic traces, called bi-traces. A bi-trace is consistent when the messages in the two symbolic traces are indistinguishable by the adversary. Informally, two protocols are open bisimilar when each action in one protocol can be simulated in the other using a consistent bi-trace.

Diff-equivalence (d) is another strictly stronger notion that is defined for protocols that have the same structure and differ only by the messages they exchange. It means that, during execution, all communications and tests, including those that the adversary can make, either succeed for both protocols or fail for both protocols. This property implies that both protocols still have the same structure during execution.

Equational theories (S). What is the support for equational theories? At a high-level, extra support for certain axioms enables detecting a larger class of attacks (see, e.g., [40], [41]). We provide a coarse classification as follows: tools that support a fixed set of equational theories or no equational theories at all (○); tools that support user-defined equational theories, but without associative-commutative (AC) axioms (○); tools that support user-defined equational theories with AC axioms (●). Supporting associative and commutative properties enables detecting a much larger class of attacks, since they allow the most detailed modeling of, e.g., xor operations, abelian groups, and Diffie-Hellman constructions. One caveat is that the finer details between these coarse classifications often make them incomparable, and even where they overlap, they are not all equally effective for analyzing concrete protocols.

Global mutable state (S). Does the tool support verification of protocols with global mutable state? Many real-world protocols involve shared databases (e.g., key servers) or shared memory, so reasoning support for analyzing complex, stateful attacks scenarios extends the reach of such tools [28].

Link to implementation (T). Can the tool extract/generate executable code from specifications in order to link symbolic security guarantees to implementations?

† **Abstractions (U).** Does the tool use abstraction? Algorithms may use abstraction to overestimate attack possibilities,

e.g., by computing a superset of the adversary’s knowledge. This can yield more efficient and fully automatic analysis systems and can be a workaround to undecidability, but comes at the cost of incompleteness, i.e., false attacks may be found or the tool may terminate with an indefinite answer.

‡ **Interactive mode (U).** Does the tool support an interactive analysis mode? Interactive modes generally trade off automation for control. While push-button tools are certainly desirable, they may fail opaquely (perhaps due to undecidability barriers), leaving it unclear or impossible to proceed. Interactive modes can allow users to analyze failed automated analysis attempts, inspect partial proofs, and to provide hints and guide analyses to overcome any barriers.

§ **Independent verifiability (T).** Are the analysis results independently machine-checkable? Symbolic tools implement complex verification algorithms and decision procedures, which may be buggy and return incorrect results. This places them in the trusted computing base. Exceptions include scyther-proof [26], which generates proof scripts that can be machine-checked in the Isabelle theorem prover [42], and SPEC [36], which can produce explicit evidence of security claims that can be checked for correctness.

Specification language (U). How are protocols specified? The categorizations are domain-specific security protocol languages (\triangleright), process calculus (\star), multiset rewriting (\ast), and general programming language (\diamond). General programming languages are arguably the most familiar to non-experts, while security protocol languages (i.e., notations for describing message flows between parties) are commonplace in cryptography. Process calculi and multiset rewriting may be familiar to formal methods practitioners. Process calculi are formal languages for describing concurrent processes and their interactions (e.g., [43]–[45]). Multiset rewriting is a more general and lower-level formalism that allows for various encodings of processes, but has no built-in notion of a process. It provides a natural formalism for complex state machines.

C. Symbolic Security: Discussion

Achievements: Symbolic proofs for real-world case studies. Of the considered symbolic tools, ProVerif and Tamarin stand out as having been used to analyze large, real-world protocols. They offer unprecedented combinations of scalability and expressivity, which enables them to deal with complex systems and properties. Moreover, they provide extensive documentation, a library of case studies, and practical usability features (e.g., packaging, a graphical user interface for Tamarin, attack reconstruction in HTML for Proverif).

Next, we provide a rough sense of their scalability on real-world case studies; more precise numbers can be found in the respective papers. It is important to keep in mind that comparisons between tools are difficult (even on similar case studies), so these numbers should be taken with a grain of salt.

ProVerif has been used to analyze TLS 1.0 [46] (seconds to several hours depending on the security property) and 1.3 [3] (around one hour), Signal [47] (a few minutes to more than a day depending on the security property), and Noise

Tool	RF	Auto	Comp	CS	Link	TCB
AutoG&P \diamond [55]	●	●	○	●	○	self, SMT
CertiCrypt $\triangleright\Diamond$ [56]	●	○	○	●	●	Coq
CryptHOL \diamond [57]	○	○	●	●	○	Isabelle
CryptoVerif $\ast\Diamond$ [58]	●	●	○	●	●	self
EasyCrypt $\triangleright\Diamond$ [59]	●	○	●	●	●	self, SMT
F7 \diamond [17]	●	○	●	○	●	self, SMT
F $\ast\Diamond$ [60]	●	○	●	○	●	self, SMT
FCF \diamond [61]	●	○	●	●	●	Coq
ZooCrypt \diamond [62]	●	●	○	●	○	self, SMT
Reasoning Focus (RF)	Concrete security (CS)			Specification language		
● – automation focus	● – security + efficiency			★ – process calculus		
○ – expressiveness focus	● – security only			\triangleright – imperative		
	○ – no support			\diamond – functional		

TABLE II
OVERVIEW OF TOOLS FOR COMPUTATIONAL SECURITY ANALYSIS. SEE SECTION II-D FOR MORE DETAILS ON COMPARISON CRITERIA.

protocols [48] (seconds to days depending on the protocol). In general, more Diffie-Hellman key agreements (e.g., in Signal and Noise) increase analysis times.

Tamarin has been used to analyze the 5G authentication key exchange protocol [49] (around five hours), TLS 1.3 [2], [4] (around one week, requiring 100GB RAM), the DNP3 SAv5 power grid protocol [50] (several minutes), and Noise protocols [51] (seconds to hours depending on the protocol).

Challenge: Verifying equivalence properties. Many security properties can be modeled accurately by equivalence properties, but they are inherently more difficult to verify than trace properties. This is because they involve relations between traces instead of single traces. As such, tool support for reasoning about equivalence properties is thus substantially less mature. For full automation, either one bounds the number of sessions or one has to use the very strong notion of diff-equivalence, which cannot handle many desired properties, e.g., vote privacy in e-voting and unlinkability.

For the bounded setting, recent developments include support for more equational theories (AKISS [32], DEEPSEC [34]), for protocols with else branches (APTE [33], AKISS, DEEPSEC) and for protocols whose actions are not entirely determined by their inputs (APTE, DEEPSEC). There have also been performance improvements based on partial order reduction (APTE, AKISS, DEEPSEC) or graph planning (SAT-Equiv). For the unbounded setting, diff-equivalence, first introduced in ProVerif [52] and later adopted by Maude-NPA [53] and Tamarin [54], remains the only fully automated approach for proving equivalences. Because trace equivalence is the most natural for formalizing privacy properties, verifying more general equivalence properties for an unbounded number of sessions remains a challenge.

D. Computational Tools: State of the Art

Table II presents a taxonomy of general-purpose computational tools. Tools are listed alphabetically and are categorized as follows.

Reasoning focus (U). Is the tool’s reasoning focus on automation (●) or on expressivity (○)? Automation focus means being able to produce automatically or with light interaction a security proof (at the cost of some expressiveness). Dually,

expressivity focus means being able to express arbitrary arguments (at the cost of some automation).

Automated proof-finding (U). Can the tool automatically find security proofs? A subset of the automation-focused tools can automatically (non-interactively) find security proofs in restricted settings (e.g., proofs of pairing-based schemes for AutoG&P, proofs of key exchange protocols using a catalog of built-in game transformations for CryptoVerif, proofs of padding-based public key encryption schemes for ZooCrypt).

Composition (U). Does the tool support compositional reasoning? Support for decomposing security arguments of cryptographic systems into security arguments for their core components is essential for scalable analysis.

Concrete security (A). Can the tool be used to prove concrete bounds on the adversary’s success probability and execution time? We consider tools with no support (○), support for success probability only (◐), and support for both (●).

Link to implementation (T). Can the tool extract/generate executable code from specifications in order to link computational security guarantees to implementations?

Trusted computing base (T). What lies in the trusted computing base (TCB)? An established general-purpose theorem prover such as Coq [63] or Isabelle [64] is usually held as the minimum TCB for proof checking. Most tools, however, rely on an implementation of the tool’s logics in a general purpose language that must be trusted (self). Automation often relies on SMT solvers [65], such as Z3 [66].

Specification language (U). What kind of specification language is used? All tools support some functional language core for expressing the semantics of operations (\diamond). Some tools support an imperative language (\triangleright) in which to write security games, while others rely on a process calculus (\star).

E. Computational Security: Discussion

Achievements: Machine-checked security for real-world cryptographic designs. Computational tools have been used to develop machine-checked security proofs for a range of real-world cryptographic mechanisms. CryptoVerif has been used for a number of protocols, including TLS 1.3 [3], Signal [47], and WireGuard [67]. EasyCrypt has been used for the Amazon Web Service (AWS) key management system [68] and the SHA-3 standard [69]. F7 was used to build miTLS, a reference implementation of TLS 1.2 with verified computational security at the code-level [70], [71]. F* was used to implement and verify the security of the TLS 1.3 record layer [1].

Takeaway: CryptoVerif is good for highly automated computational analysis of protocols and systems. CryptoVerif is both a proof-finding and proof-checking tool. It works particularly well for protocols (e.g., key exchange), as it can produce automatically or with a light guidance a sequence of proof steps that establish security. One distinctive strength of CryptoVerif is its input language based on the applied π -calculus [45], which is well-suited to describing protocols that exchange messages in sequence. Another strength of CryptoVerif is a carefully crafted modeling of security assumptions that help the automated discovery of proof steps. In turn,

automation is instrumental to deal with large cryptographic games and games that contain many different cases, as is often the case in proofs of protocols.

Takeaway: F* is good for analysis of full protocols and systems. F* is a general-purpose verification-oriented programming language. It works particularly well for analyzing cryptographic protocols and systems beyond their cryptographic core. Computational proofs in F* rely on transforming a detailed protocol description into a final (ideal) program by relying on ideal functionalities for cryptographic primitives. Formal validation of this transformation is carried out manually, with some help from the F* verification infrastructure. Formal verification of the final program is done within F*. This approach is driven by the insight that critical security issues, and therefore also potential attacks, often arise only in detailed descriptions of full protocols and systems (compared to when reasoning about cryptographic cores). The depth of this insight is reflected by the success of F*-based verification both in helping discovering new attacks on real-world protocols like TLS [72], [73] as well as in verifying their concrete design and implementation [1], [70].

Takeaway: EasyCrypt is the closest to pen-and-paper cryptographic proofs. EasyCrypt supports a general-purpose relational program logic (i.e., a formalism for specifying and verifying properties about two programs or two runs of the same program) that captures many of the common game hopping techniques. This is complemented by libraries that support other common techniques, e.g., the PRF/PRP switching lemma, hybrid arguments, and lazy sampling [8]. In addition, EasyCrypt features a union bound logic for upper bounding the probability of some event E in an experiment (game) G (e.g., bounding the probability of collisions in experiments that involve hash functions). Overall, EasyCrypt proofs closely follow the structure of pen-and-paper arguments. A consequence is that EasyCrypt is amenable to proving the security of primitives, as well as protocols and systems.

Challenge: Scaling security proofs for cryptographic systems. Analyzing large cryptographic systems is best done in a modular way by composing simpler building blocks. However, cryptographers have long recognized the difficulties of preserving security under composition [74]. Most game-based security definitions do not provide out-of-the-box composition guarantees, so simulation-based definitions are the preferred choice for analyzing large cryptographic systems, with universal composability (UC) being the gold-standard—UC definitions guarantee secure composition in arbitrary contexts [9]. Work on developing machine-checked UC proofs is relatively nascent [75]–[77], but is an important and natural next step for computational tools.

F. Further Reading

Another class of tools leverages the benefits of automated verification to support automated synthesis of secure cryptographic designs, mainly in the computational world [62], [78]–[81]. Cryptographic compilers provide high-level abstractions (e.g., a domain-specific language) for describing cryptographic

tasks, which are then compiled into custom protocol implementations. These have been proposed for verifiable computation [82]–[85], zero-knowledge [86]–[89], and secure multiparty computation [90] protocols, which are parameterized by a proof-goal or a functionality to compute. Some are supported by proofs that guarantee the output protocols are correct and/or secure for every input specification [91]–[94]. We recommend readers to also consult other related surveys. Blanchet [95] surveys design-level security until 2012 (with a focus on ProVerif). Cortier et al. [96] survey computational soundness results, which transfer security properties from the symbolic world to the computational world.

III. FUNCTIONAL CORRECTNESS AND EFFICIENCY

In this section, we focus on the role of computer-aided cryptography in developing functionally correct and efficient implementations.

A. Critical Review

Why are functional correctness and efficiency important?

To reap the benefits of design-level security guarantees, implementations must be an accurate translation of the design proven secure. That is, they must be functionally correct (i.e., have equivalent input/output behavior) with respect to the design specification. Moreover, to meet practical deployment requirements, implementations must be efficient. Cryptographic routines are often on the critical path for security applications (e.g., for reading and writing TLS packets or files in an encrypted file system), and so even a few additional clock-cycles can have a detrimental impact on overall performance.

How can functional correctness and efficiency fail?

If performance is not an important goal, then achieving functional correctness is relatively easy—just use a reference implementation that does not deviate too far from the specification, so that correctness is straightforward to argue. However, performance demands drive cryptographic code into extreme contortions that make functional correctness difficult to achieve, let alone prove. For example, OpenSSL is one of the fastest open source cryptographic libraries; they achieve this speed in part through the use of Perl code to generate strings of text that additional Perl scripts interpret to produce input to the C preprocessor, which ultimately produces highly tuned, platform-specific assembly code [103]. Many more examples of high-speed crypto code written at assembly and pre-assembly levels can be found in SUPERCOP [107], a benchmarking framework for cryptography implementations.

More broadly, efficiency considerations typically rule out exclusively using high-level languages. Instead, C and assembly are the de facto tools of the trade, adding memory safety to the list of important requirements. Indeed, memory errors can compromise secrets held in memory, e.g., in the Heartbleed attack [108]. Fortunately, as we discuss below, proving memory safety is table stakes for most of the tools we discuss. Additionally, achieving best-in-class performance demands aggressive, platform-specific optimizations, far beyond what is achievable by modern optimizing compilers (which are

problematic in their own ways, as we will see in Section IV). Currently, these painstaking efforts are manually repeated for each target architecture.

How are these failures being addressed outside CAC?

Given its difficulty, the task of developing high-speed cryptography is currently entrusted to a handful of experts. Even so, experts make mistakes (e.g., a performance optimization to OpenSSL’s AES-GCM implementation nearly reached deployment even though it enabled arbitrary message forgeries [109]; an arithmetic bug in OpenSSL led to a full key recovery attack [110]). Current solutions for preventing more mistakes are (1) auditing, which is costly in both time and expertise, and (2) testing, which cannot be complete for the size of inputs used in cryptographic algorithms. These solutions are also clearly inadequate: Despite widespread usage and scrutiny, OpenSSL’s `libcrypto` library reported 24 vulnerabilities between January 1, 2016 and May 1, 2019 [7].

How can computer-aided cryptography help? Cryptographic code is an ideal target for program verification. Such code is both critically important and difficult to get right. The use of heavyweight formal methods is perhaps the only way to attain the high-assurance guarantees expected of them. At the same time, because the volume of code in cryptographic libraries is relatively small (compared to, say, an operating system), verifying complex, optimized code is well within reach of existing tools and reasonable human effort, without compromising efficiency.

What are the fine-print caveats? Functional correctness makes implicit assumptions, e.g., correct modeling of hardware functional behavior. Another source of implicit assumptions is the gap between code and verified artifacts, e.g., verification may be carried out on a verification-friendly representation of the source code, rather than on the source code itself. Moreover, proofs may presuppose correctness of libraries, e.g., for efficient arithmetic. Finally, as with any software, verification tools may have bugs.

What background do I need to know? Functional correctness is the central focus of program verification. An implementation can be proved functionally correct in two different ways: equivalence to a reference implementation, or satisfying a functional specification, typically expressed as pre-conditions (what the program requires on inputs) and post-conditions (what the program guarantees on outputs). Both forms of verification are supported by a broad range of tools. A unique aspect of cryptographic implementations is that their correctness proofs often rest on non-trivial mathematics. Mechanizing them thus requires striking a good balance between automation and user control. Nevertheless, SMT-based automation remains instrumental for minimizing verification effort, and almost all tools offer an SMT-based backend.

Typically, functional correctness proofs are carried out at source level. A long-standing challenge is how to carry guarantees to machine code. This can be addressed using verified compilers, which are supported by formal correctness proofs. CompCert [111] is a prime example of moderately optimizing verified compiler for a large fragment of C. However, the

Tool	Memory safety	Automation	Parametric verification	Input language	Target(s)	TCB
Cryptol + SAW [97]	●	●	○	C, Java	C, Java	SAT, SMT
CryptoLine [98]	○	●	○	CryptoLine	C	Boolector, MathSAT, Singular
Dafny [99]	●	●	○	Dafny	C#, Java, JavaScript, Go	Boogie, Z3
F* [60]	●	●	○	F*	OCaml, F#, C, Asm, Wasm	Z3, typechecker
Fiat Crypto [6]	●	○	●	Gallina	C	Coq, C compiler
Frama-C [100]	●	●	○	C	C	Coq, Alt-Ergo, Why3
gfvcrif [101]	○	●	○	C	C	g++, Sage
Jasmin [102]	●	●	○	Jasmin	Asm	Coq, Dafny, Z3
Vale [103], [104]	●	●	●	Vale	Asm	Dafny or F*, Z3
VST [105]	●	○	○	Gallina	C	Coq
Why3 [106]	○	●	○	WhyML	OCaml	SMT, Coq

Automation
 ● – automated ● – automated + interactive ○ – interactive

TABLE III

OVERVIEW OF TOOLS FOR FUNCTIONAL CORRECTNESS. SEE SECTION III-B FOR MORE DETAILS ON COMPARISON CRITERIA.

trade-off is that verified compilers typically come with fewer optimizations than mainstream compilers and target fewer platforms.

B. Program Verification Tools: State of the Art

Table III presents a taxonomy of program verification tools that have been used for cryptographic implementation. Tools are listed alphabetically and are categorized as follows.

Memory-safety (S). Can the tool verify that programs are memory safe? Memory safety ensures that all runs of a program are free from memory errors (e.g., buffer overflow, null pointer dereferences, use after free).

Automation (U). Tools provide varying levels of automation. We give a coarse classification: automatic tools (●), tools that combine automated and interactive theorem proving (●), and tools that allow only interactive theorem proving (○).

Parametric verification (U). Can the tool verify parameterized implementations? This enables writing and verifying generic code that can be used to produce different implementations depending on the supplied parameters. For example, Fiat Crypto [6] can generate verified elliptic curves implementations parameterized by a prime modulus, limb representation of field elements, and hardware platform; Vale [103], [104] implementations are parameterized by the operating system, assembler, and hardware platform.

Input language (U). What is the input language? Many toolchains use custom verification-oriented languages. Dafny is a high-level imperative language, whereas F*, Gallina (used in Coq), and WhyML (used in Why3) are functional languages. CryptoLine, Jasmin, and Vale are assembly-like languages; Jasmin and Vale provide high-level control-flow structures such as procedures, conditionals, and loops. Other tools take code written in existing languages (e.g., C, Java).

Target(s) (A,S). At what level is the analysis carried out (e.g., source-level or assembly-level)? Note that tools targeting source-level analysis must use verified compilers (e.g., CompCert [111]) to carry guarantees to machine-level, which comes with a performance penalty. Tools targeting assembly-level analysis sidestep this dilemma, but generally verification becomes more difficult.

Trusted computing base (T). What lies in the trusted computing base? Many verification frameworks rely on building-

Implementation	FC	CT	Tool(s)	Target	% faster
evercrypt [7]	●	●	F*, Vale	64-bit C, Intel ADX asm	25.92
precomp [112]	○	○	—	Intel ADX asm	25.77
sandy2x [113]	○	○	—	Intel AVX asm	11.15
hac1 [7]	●	○	F*	64-bit C	8.69
jasmin [102]	○	●	Jasmin	Intel x86_64 asm	7.88
amd64 [114]	●	○	Coq, SMT	Intel x86_64 asm	6.11
fiat [6]	●	○	Fiat Crypto	64-bit C	5.39
donna64 [115]	○	○	—	64-bit C	0.00

Functional correctness (FC), Constant-time (CT)
 ● – verified ● – partially verified ○ – not verified

TABLE IV

COMPARISON OF CURVE25519 IMPLEMENTATIONS. % FASTER CALCULATED USING DONNA64 AS THE BASELINE.

block verification tools, such as SMT solvers (e.g., Z3) and interactive theorem provers (e.g., Coq). While these are acknowledged to be important trust assumptions of verification tools, verified artifacts tend to rely on additional trust assumptions, e.g., unverified interoperability between tools or only verifying small routines in a larger primitive.

C. Discussion

Achievements: Verified primitives are being deployed at Internet-scale. A recent milestone achievement of computer-aided cryptography is that verified primitives are being deployed at scale. Verified primitives in the HAC1* [5] library are used in Mozilla Firefox’s NSS security engine, and verified elliptic curve implementations in the Fiat Cryptography library [6] are used in Google’s BoringSSL library.

There are several common insights to these successes. First, verified code needs to be as fast or faster than the code being replaced. Second, verified code needs to fit the APIs that are actually in use. Third, it helps if team members work with or take internships with the companies that use the code. In the case of HAC1*, it additionally helped that they replaced an entire ciphersuite, and that they were willing to undertake a significant amount of non-research work, such as packaging and testing, that many academic projects stop short of.

Takeaway: Verified implementations are now as fast or faster than their unverified counterparts. Through decades of research in formal verification, it was commonly accepted that the proof burden in verifying complex, optimized code was exorbitant; verified code would be hard-pressed to com-

pete with unverified code in terms of performance. However, various projects in the cryptography domain have challenged this position. We are seeing verified implementations that meet the performance of the fastest unverified implementations. We conclude that there is currently no conceptual or technological barrier that prevents verifying the fastest implementations available, although more effort is expected.

As a small case study, we look at Curve25519 [116], a widely used elliptic curve that has received considerable interest from the applied cryptography community (in setting new speed records) and the formal methods community (in verifying that high-speed implementations are correct and secure). We compare a number of Curve25519 implementations in Table IV. These comprise some of the fastest available verified and unverified implementations; they are written in C, assembly, or a combination of both.

To compare their performance, we measure the number of CPU cycles (median over 5K executions) it takes to perform scalar multiplication. We report the performance increase (% faster) over donna64 [115], one of the fastest known (unverified) C implementations. All measurements are collected on a 1.8 GHz Intel i7-8565U CPU with 16 GB of RAM; hyper-threading and dynamic-processor scaling (e.g., Turbo Boost) are disabled. Implementations written in C are compiled using GCC 9.2 with optimization flag -O3. To summarize, several verified C implementations (hacl and fiat) beat donna64; the fastest verified assembly implementation (evercrypt) meets the fastest unverified assembly implementation (precomp).

Takeaway: Higher performance entails larger verification effort. Verifying generic, high-level code is typically easier, but comes with a performance cost. Hand-written assembly can achieve best in class performance by taking advantage of hardware-specific optimizations, but verifying such implementations is quite difficult due to complex side-effects, unstructured control-flow, and flat structure. Moreover, this effort must be repeated for each platform. C code is less efficient, as hardware-specific features are not a part of standard portable C, but implementations need only be verified once and can then be run on any platform. Code written in higher-level languages is even less efficient, but verification becomes much easier (e.g., memory safety can be obtained for free). These aspects are discussed further in the Vale and Jasmin papers [102], [103], [117].

Challenge: Automating equivalence proofs. Significant progress could be made if functional correctness proofs could be solved by providing a sequence of simple transformations that connect specifications to targets and relying on an automatic tool to check these simple transformations. Promising recent work in this direction [118] demonstrates the feasibility of the approach. However, the current approaches are not automatic: neither in finding the transformations nor in proving them. The latter seems achievable for many useful control-flow-preserving transformations, whereas the former could be feasible at least for common control-flow transformations.

Challenge: Functional correctness of common arithmetic routines. Verifying cryptographic code often involves tricky

mathematical reasoning that SMT-based tools can struggle with. Examples range from proving the correctness of the Montgomery representations [119] used to accelerate big-integer computations, to the nuts-and-bolts of converting between, say, 64-bit words and the underlying bytes. At present, most verification efforts build this infrastructure from scratch and customize it for their own particular needs, which leads to significant duplication of effort across projects. Hence, an open challenge is to devise a common core of such routines (e.g., a verified version of the GMP library [120]) that can be shared across all (or most) verification projects, despite their reliance on different tools and methodologies.

D. Further Reading

While our principal focus is on cryptographic code, verifying systems code is an important and active area of research. For example, there has been significant work in verifying operating systems code [121]–[127], distributed systems [128]–[130], and even entire software stacks [131]. We expect that these two strands of work will cross paths in the future.

IV. IMPLEMENTATION-LEVEL SECURITY

In this section, we focus on the role of computer-aided cryptography in establishing implementation-level security guarantees, with a particular focus on software protections against digital side-channel attacks. Hardware protections are beyond the scope of this paper and are left as further reading. By digital side-channel attacks, we mean those that can be launched by observing intentionally exposed interfaces by the computing platform, including all execution time variations and observable side-effects in shared resources such as the cache. This excludes physical side channels such as power consumption, electromagnetic radiation, etc.

A. Critical Review

Why is implementation-level security important? Although design-level security can rule out large classes of attacks, guarantees are proven in a model that idealizes an attacker’s interface with the underlying algorithms: They can choose inputs and observe outputs. However, in practice, attackers can observe much more than just the functional behavior of cryptographic algorithms. For example, side-channels are interfaces available at the implementation-level (but unaccounted for at the design-level) from which information can leak as side-effects of the computation process (e.g., timing behavior, memory access patterns). And indeed, these sources of leakage are devastating—key-recovery attacks have been demonstrated on real implementations, e.g., on RSA [142] and AES [143].

How can implementation-level security fail? The prevailing technique for protecting against digital side-channel attacks is to follow *constant-time* coding guidelines [144]. We stress that the term is a bit of a misnomer: The idea of constant-time is that an implementation’s logical execution time (not wall-clock execution time) should be independent of the values of secret data; it may, however, depend on public data, such as input length. To achieve this, constant-time implementations must

Tool		Target	Method	Synthesis	Sound	Complete	Public inputs	Public outputs	Control flow	Memory access	Variable-time op.
ABPV13	[132]	C	DV	○	●	●	●	○	●	●	○
CacheAudit	[133]	Binary	Q	○	○	○	○	○	○	○	○
ct-verif	[134]	LLVM	DV	○	●	●	●	●	●	●	●
CT-Wasm	[135]	Wasm	TC	○	●	○	●	○	●	●	●
FaCT	[136]	LLVM	TC	●	●	○	●	○	●	●	●
FlowTracker	[137]	LLVM	DF	○	●	○	●	○	●	●	○
Jasmin	[102]	asm	DV	○	●	●	●	●	●	●	○
KMO12	[138]	Binary	Q	○	●	○	○	○	○	○	○
Low*	[139]	C	TC	○	●	○	●	○	●	●	●
SC Eliminator	[140]	LLVM	DF	●	●	○	●	○	●	●	○
Vale	[103]	asm	DF	○	●	○	●	●	●	●	●
VirtualCert	[141]	x86	DF	○	●	○	●	○	●	●	○

Method
 TC – type-checking DF – data-flow analysis DV – deductive verification Q – Quantitative

TABLE V

OVERVIEW OF TOOLS FOR SIDE-CHANNEL RESISTANCE. SEE SECTION IV-B FOR MORE DETAILS ON TOOL FEATURES.

follow a number of strict guidelines, e.g., they must avoid variable-time operations, control flow, and memory access patterns that depend on secret data. Unfortunately, complying with constant-time coding guidelines forces implementors to avoid natural but potentially insecure programming patterns, making enforcement error-prone.

Even worse, the observable properties of a program’s execution are generally not evident from source code alone. Thus, software-invisible optimizations, e.g., compiler optimizations or data-dependent instruction set architecture (ISA) optimizations, can remove source-level countermeasures. Programmers also assume that the computing machine provides memory isolation, which is a strong and often unrealistic assumption in general-purpose hardware (e.g., due to isolation breaches allowed by speculative execution mechanisms).

How are these failures being addressed outside CAC?

To check that implementations correctly adhere to constant-time coding guidelines, current solutions are (1) auditing, which is costly in both time and expertise, and (2) testing, which commits the fallacy of interpreting constant-time to be constant wall-clock time. These solutions are inadequate: A botched patch for a timing vulnerability in TLS [145] led to the Lucky 13 timing vulnerability in OpenSSL [146]; in turn, the Lucky 13 patch led to yet another timing vulnerability [147]!

To prevent compiler optimizations from interfering with constant-time recipes applied at the source-code level, implementors simply avoid using compilers at all, instead choosing to implement cryptographic routines and constant-time recipes directly in assembly. Again, checking that countermeasures are implemented correctly is done through auditing and testing, but in a much more difficult, low-level setting.

Dealing with micro-architectural attacks that breach memory isolation, such as Spectre and Meltdown [148], [149], is still an open problem and seems to be out of reach of purely software-based countermeasures if there is to be any hope of achieving decent performance.

How can computer-aided cryptography help? Program analysis and verification tools can automatically (or semi-automatically) check whether a given implementation meets constant-time coding guidelines, thereby providing a formal foundation supporting heretofore informal best practices. Even

further, some tools can automatically repair code that violates constant-time into compliant code. These approaches necessarily abstract the leakage interface available to real-world attackers, but being precisely defined, they help clarify the gap between formal leakage models and real-world leakage.

What are the fine-print caveats? Implementation-level proofs are only as good as their models, e.g., of physically observable effects of hardware. Furthermore, new attacks may challenge these models. Implicit assumptions arise from gaps between code and verified artifacts.

What background do I need to know? Formal reasoning about side-channels is based on a *leakage model*. This model is defined over the semantics of the target language, abstractly representing what an attacker can observe during the computation process. For example, the leakage model for a branching operation may leak all program values associated with the branching condition. After having defined the appropriate leakage models, proving that an implementation is secure (with respect to the leakage models) amounts to showing that the leakage accumulated over the course of execution is independent of the values of secret data. This property is an instance of *observational non-interference*, an information flow property requiring that variations in secret data cause no differences in observable outputs [150].

The simplest leakage model is the program counter policy, where the program control-flow is leaked during execution [151]. The most common leakage model, namely the constant-time policy, additionally assumes that memory accesses are leaked during execution. This leakage model is usually taken as the best practice to remove exploitable execution time variations and a best-effort against cache-attacks launched by co-located processes. A more precise leakage model called the size-respecting policy also assumes that operand sizes are leaked for specific variable-time operations. For more information on leakage models, see the paper by Barthe et al. [150, Section IV.D].

B. Digital Side-Channel Tools: State of the Art

Table V presents a taxonomy of tools for verifying digital side-channel resistance. Tools are listed alphabetically and are categorized as follows.

Target (A,S). At what level is the analysis performed (e.g., source, assembly, binary)? To achieve the most reliable guarantees, analysis should be performed as close as possible to the executed machine code.

Method (A). The tools we consider all provide a means to verify absence of timing leaks in a well-defined leakage model, but using different techniques:

- Static analysis techniques use type systems or data-flow analysis to keep track of data dependencies from secret inputs to problematic operations.
- Quantitative analysis techniques construct a rich model of a hardware feature, e.g. the cache, and derive an upper-bound on the leaked information.
- Deductive verification techniques prove that the leakage traces of two executions of the program coincide if the public parts of the inputs match. These techniques are closely related to the techniques used for functional correctness.

Type-checking and data-flow analysis are more amenable to automation, and they guarantee non-interference by excluding all programs that could pass secret information to an operation that appears in the trace. The emphasis on automation, however, limits the precision of the techniques, which means that secure programs may be rejected by the tools (i.e., they are not complete). Tools based on deductive verification are usually complete, but require more user interaction. In some cases, users interact with the tool by annotating code, and in others the users use an interactive proof assistant to complete the proof. It is hard to conciliate a quantitative bound on leakage with standard cryptographic security notions, but such tools can also be used to prove a zero-leakage upper bound, which implies non-interference in the corresponding leakage model.

Synthesis (U). Can the tool take an insecure program and automatically generate a secure program? Tools that support synthesis (e.g., FaCT [136] and SC Eliminator [140]) can automatically generate secure implementations from insecure implementations. This allows developers to write code naturally with constant-time coding recipes applied automatically.

Soundness (A, T). Is the analysis sound, i.e., it only deems secure programs as secure? Note that this is our baseline filter for consideration, but we make this explicit in the table. Still, it bears mentioning that some unsound tools are used in practice. One example is ctgrind [152], an extension of Valgrind that takes in a binary with taint annotations and checks for constant-address security via dynamic analysis. It supports public inputs but not public outputs, and is neither sound nor complete.

Completeness (A, S). Is the analysis complete, i.e., it only deems insecure programs as insecure?

Public input (S). Does the tool support public inputs? Support for public inputs allows differentiating between public and secret inputs. Implementations can benignly violate constant-time policies without introducing side-channel vulnerabilities by leaking no more information than public inputs of computations. Unfortunately, tools without such support would reject these implementations as insecure; forcing execution behaviors

to be fully input independent may lead to large performance overheads.

Public output (S). Does the tool support public outputs? Similarly, support for public outputs allows differentiating between public and secret outputs. The advantages to supporting public outputs is the same as those for supporting public inputs: for example, branching on a bit that is revealed to the attacker explicitly is fine.

Control flow leakage (S). Does the tool consider control-flow leakage? The leakage model includes values associated with conditional branching (e.g., if, switch, while, for statements) during program execution.

Memory access leakage (S). Does the tool consider memory access pattern leakage? The leakage model includes memory addresses accessed during program execution.

Variable-time operation leakage (S). Does the tool consider variable-time operation leakage? The leakage model includes inputs to variable-time operations (e.g., floating point operations [153]–[155], division and modulus operations on some architectures) classified according to timing-equivalent ranges.

C. Discussion

Achievements: Automatic verification of constant-time real-world code. There are several tools that can perform verification of constant-time code automatically, both for high-level code and low-level code. These tools have been applied to real-world libraries. For example, portions of the assembly code in OpenSSL have been verified using Vale [103], high-speed implementations of SHA-3 and TLS 1.3 ciphersuites have been verified using Jasmin [102], and various off-the-shelf libraries have been analyzed with FlowTracker [137].

Takeaway: Lowering the target provides better guarantees. Of the surveyed tools, several operate at the level of C code; others operate at the level of LLVM assembly; still others operate at the level of assembly or binary. The choice of target is important. To obtain a faithful correspondence with the executable program under an attacker’s scrutiny, analysis should be performed as close as possible to the executed machine code. Given that mainstream compilers (e.g., GCC and Clang) are known to optimize away defensive code and even introduce new side-channels [156], compiler optimizations can interfere with countermeasures deployed and verified at source-level.

Challenge: Secure, constant-time preserving compilation. Given that mainstream compilers can interfere with side-channel countermeasures, many cryptography engineers avoid using compilers at all, instead choosing to implement cryptographic routines directly in assembly, which means giving up the benefits of high-level languages.

An alternative solution is to use secure compilers that carry source-level countermeasures along the compilation chain down to machine code. This way, side-channel resistant code can be written using portable C, and the secure compiler takes care of preserving side-channel resistance to specific architectures. Barthe et al. [150] laid the theoretical foundations of constant-time preserving compilation. These ideas were subsequently realized in the verified CompCert C compiler [157].

Unfortunately, CompCert-generated assembly code is not as efficient as that generated by GCC and Clang, which in turn lags the performance of hand-optimized assembly.

Challenge: Protecting against micro-architectural attacks.

The constant-time policy is designed to capture logical timing side channels in a simple model of hardware. Unfortunately, this simple model is inappropriate for modern hardware, as microarchitectural features, e.g., speculative or out-of-order execution, can be used for launching devastating side-channel attacks. Over the last year, the security world has been shaken by a series of attacks, including Spectre [148] and Meltdown [149]. A pressing challenge is to develop notions of constant-time security and associated verification methods that account for microarchitectural features.

Challenge: Rethinking the hardware-software contract from secure, formal foundations. An ISA describes (usually informally) what one needs to know to write a functionally correct program [158], [159]. However, current ISAs are an insufficient specification of the hardware-software contract when it comes to writing secure programs [160]. They do not capture hardware features that affect the temporal behavior of programs, which makes carrying side-channel countermeasures at the software-level to the hardware-level difficult.

To rectify this, researchers have called on new ISA designs that expose, for example, the temporal behaviors of hardware, which can lend to reasoning about them in software [160]. This, of course, poses challenging and competing requirements for hardware architects, but we believe developing formal foundations for verification and reasoning about security at the hardware-software interface can help. This line of work seems also to be the only path that can lead to a sound, formal treatment of micro-architectural attacks.

D. Further Reading

For lack of space, we had to omit several threads of relevant work, e.g., on verifying side-channel resistance in hardware [161]–[165], and on verifying masked implementations aimed at protecting against differential power analysis attacks [166]–[171].

V. CASE STUDY I: CONSOLIDATING GUARANTEES

Previous sections focus on specific guarantees: design-level security, functional correctness, efficiency, and side-channel resistance. This case study focuses on unifying approaches that can combine these guarantees. This is a natural and important step towards the Holy Grail of computer-aided cryptography: to deliver guarantees on executable code that match the strength and elegance of guarantees on cryptographic designs.

Table VI collects implementations that verifiably meet more than one guarantee. Implementations are grouped by year (demarcated by dashed lines), starting from 2014 and ending in 2019; within each year, implementations are listed alphabetically by author. We report on the primitives included, the languages targeted, the tools used, and the guarantees met.

Computational security. We categorize computational security guarantees as follows: verified (●), partially verified (◐), not verified (○), and not applicable (–). The

HACL*-related implementations are partially verified, as only AEAD primitives have computational proofs, which are semi-mechanized [1]. Security guarantees do not apply to, e.g., elliptic curve implementations or bigint code.

Functional correctness. We categorize functional correctness guarantees as follows: target-level (●), source-level (◐), and not verified (○). Target-level guarantees can be achieved in two ways: Either guarantees are established directly on assembly code, or guarantees are established at source level and a verified compiler is used.

Efficiency. We categorize efficiency as follows: comparable to assembly reference implementations (●), comparable to portable C reference implementations (◐), and slower than portable C reference implementations (○).

Side-channel resistance. We categorize side-channel resistance guarantees as follows: target-level (●), source-level (◐), and not verified (○).

Takeaway: Existing tools can be used to achieve the “grand slam” of guarantees for complex cryptographic primitives. Ideally, we would like computational security guarantees, (target-level) functional correctness, efficiency, and (target-level) side-channel guarantees to be connected in a formal, machine-checkable way (the “grand slam” of guarantees). Many implementations come close, but so far, only one meets all four. Almeida et al. [69] formally verify an efficient implementation of the sponge construction from the SHA-3 standard. It connects proofs of random oracle (RO) indistinguishability for a pseudo-code description of the sponge construction, and proofs of functional correctness and side-channel resistance for an efficient, vectorized, implementation. The proofs are constructed using EasyCrypt and Jasmin. Other works focus on either provable security or efficiency, plus functional correctness and side-channel resistance. This disconnect is somewhat expected. Provable security guarantees are established for pseudo-code descriptions of constructions, whereas efficiency considerations demand non-trivial optimizations at the level of C or assembly.

Takeaway: Integration can deliver strong and intuitive guarantees. Interpreting verification results that cover multiple requirements can be very challenging, especially because they may involve (all at once) designs, reference implementations, and optimized assembly implementations. To simplify their interpretation, Almeida et al. [174] provide a modular methodology to connect the different verification efforts, in the form of an informal meta-theorem, which concludes that an optimized assembly implementation is secure against implementation-level adversaries with side-channel capabilities. The meta-theorem states four conditions: (i) the design must be provably black-box secure in the (standard) computational model; (ii) the design is correctly implemented by a reference implementation; (iii) the reference implementation is functionally equivalent to the optimized implementation; (iv) the optimized implementation is protected against side-channels. These conditions yield a clear separation of concerns, which reflects the division of the previous sections.

Takeaway: Achieving broad scope and efficiency. Many

Implementation(s)	Target(s)	Tool(s) used	Computational security	Functional correctness	Efficiency	Side-channel resistance
RSA-OEAP [172]	C	EasyCrypt, Frama-C, CompCert	●	●	○	●
Curve25519 scalar mult. loop [114]	asm	Coq, SMT	—	●	●	○
SHA-1, SHA-2, HMAC, RSA [131]	asm	Dafny, BoogieX86	—	●	●	○
HMAC-SHA-2 [173]	C	FCF, VST, CompCert	●	●	○	○
MBE-CBC [174]	C	EasyCrypt, Frama-C, CompCert	●	●	○	●
Salsa20, AES, ZUC, FFS, ECDSA, SHA-3 [175]	Java, C	Cryptol, SAW	○	○	○	○
Curve25519 [176]	OCaml	F*, Sage	—	○	○	○
Salsa20, Curve25519, Ed25519 [102]	asm	Jasmin	○	○	●	●
SHA-2, Poly1305, AES-CBC [103]	asm	Vale	○	●	○	●
HMAC-DRBG [177]	C	FCF, VST, CompCert	●	●	○	○
HACL* ¹ [5]	C	F*	○	○	○	○
HACL* ¹ [5]	C	F*, CompCert	○	●	○	●
HMAC-DRBG [178]	C	Cryptol, SAW	○	○	○	○
SHA-3 [69]	asm	EasyCrypt, Jasmin	●	●	●	●
ChaCha20, Poly1305 [117]	asm	EasyCrypt, Jasmin	○	●	●	●
BGW multi-party computation protocol [179]	OCaml	EasyCrypt, Why3	●	○	○	○
Curve25519, P-256 [6]	C	Fiat Crypto	—	○	○	○
Poly1305, AES-GCM [104]	asm	F*, Vale	○	●	●	●
Bignum code ⁴ [98]	C	CryptoLine	—	●	○	○
WHACL* ¹ , LibSignal* [180]	Wasm	F*	○	●	○	●
EverCrypt ² [7]	C	F*	○	○	○	○
EverCrypt ³ [7]	asm	F*, Vale	○	●	●	●

Computational security	Functional correctness	Efficiency	Side-channel resistance
● – verified	● – target-level	● – comparable to asm ref	● – target-level
○ – partially verified	○ – source-level	○ – comparable to C ref	○ – source-level
○ – not verified	○ – not verified	○ – slower than C ref	○ – not verified
— – not applicable			

¹(ChaCha20, Salsa20, Poly1305, SHA-2, HMAC, Curve25519, Ed25519)

²(MD5, SHA-1, SHA-2, HMAC, Poly1305, HKDF, Curve25519, ChaCha20)

³(AES-GCM, ChaCha20, Poly1305, SHA-2, HMAC, HKDF, Curve25519, Ed25519, P-256) ⁴(In NaCl, wolfSSL, OpenSSL, BoringSSL, Bitcoin)

TABLE VI

VERIFIED CRYPTOGRAPHIC IMPLEMENTATIONS AND THEIR FORMAL GUARANTEES.

implementations target *either* C *or* assembly. This involves trade-offs between the portability and lighter verification-effort of C code, and the efficiency that can be gained via hand-tuned assembly. EverCrypt [7] is one of the first systems to target both. This garners the advantages of both, and it helps explain, in part, the broad scope of algorithms EverCrypt covers. Generic functionality and outer loops can be efficiently written and verified in C, whereas performance-critical cores can be verified in assembly. Soundly mixing C and assembly requires careful modeling of interoperation between the two, including platform and compiler-specific calling conventions, and differences in the “natural” memory and leakage models used to verify C versus assembly [7], [104].

VI. CASE STUDY II: LESSONS LEARNED FROM TLS

The Transport Layer Security (TLS) protocol is widely used to establish secure channels on the Internet, and is arguably the most important real-world deployment of cryptography to date. Before TLS version 1.3, the protocol’s design phases did not involve substantial academic analysis, and the process was highly reactive: When an attack was found, interim patches would be released for the mainstream TLS libraries or a longer-term fix would be incorporated in the next version of the standard. This resulted in an endless cycle of attacks and patches. Given the complexity of the protocol, early academic analyses considered only highly simplified cryptographic cores. However, once the academic community started considering more detailed aspects of the protocol, many new attacks were discovered, e.g., [181], [182].

The situation changed substantially during the proactive design process of TLS version 1.3: The academic community was actively consulted and encouraged to provide analysis during the process of developing multiple drafts. (See [183] for a more detailed account of TLS’s standardization history.)

On the computer-aided cryptography side of things, there were substantial efforts in verifying implementations of TLS 1.3 [1], [3] and using tools to analyze symbolic [2]–[4] and computational [3] models of TLS. Below we collect the most important lessons learned from TLS throughout the years.

Lesson: The process of formally specifying and verifying a protocol can reveal flaws. The work surrounding TLS has shown that the process of formally verifying TLS, and perhaps even just formally specifying it, can reveal flaws. The implementation of TLS 1.2 with verified cryptographic security by Bhargavan et al. [70] discovered new alert fragmentation and fingerprinting attacks and led to the discovery of the Triple Handshake attacks [72]. The symbolic analysis of TLS 1.3 draft 10 using Tamarin by Cremers et al. [2] uncovered a potential attack allowing an adversary to impersonate a client during a PSK-resumption handshake, which was fixed in draft 11. The symbolic analysis of TLS 1.3 using ProVerif by Bhargavan et al. [3] uncovered a new attack on 0-RTT client authentication that was fixed in draft 13. The symbolic analysis of draft 21 using Tamarin by Cremers et al. [4] revealed unexpected behavior that inhibited certain strong authentication guarantees. In nearly all cases, these discoveries led to improvements to the protocol, and otherwise clarified documentation of security guarantees.

Lesson: Cryptographic protocol designs are moving targets; machine-checked proofs can be more easily updated.

The TLS 1.3 specification was a rapidly moving target, with significant changes being effected on a fairly regular basis. As changes were made between a total of 28 drafts, previous analyses were often rendered stale within the space of a few months, requiring new analyses and proofs. An important benefit of machine-checked analyses and proofs over their manual counterparts is that they can be more easily and reliably updated from draft to draft as the protocol evolves [2]–[4]. Moreover, machine-checked analyses and proofs can ensure that new flaws are not introduced as components are changed.

Lesson: Standardization processes can facilitate analysis by embracing minor changes that simplify security arguments and help modular reasoning. In contrast to other protocol standards, the TLS 1.3 design incorporates many suggestions from the academic community. In addition to security fixes, these include changes purposed to simplify security proofs and automated analysis. For example, this includes changes to the key schedule that help with key separation, thus simplifying modular proofs; a consistent tagging scheme; and including more transcript information in exchanges, which simplifies consistency proofs. These changes have negligible impact on the performance of the protocol, and have helped make analyzing such a complex protocol feasible.

VII. CONCLUDING REMARKS

A. Recommendations to Authors

Our first recommendation concerns the clarity of trust assumptions. We observe that, in some papers, the distinction between what parts of an artifact are trusted/untrusted is not always clear, which runs the risk of hazy/exaggerated claims. On one hand, crisply delineating between what is trusted/untrusted may be difficult, especially when multiple tools are used, and authors may be reluctant to spell out an artifact’s weaknesses. On the other hand, transparency and clarity of trust assumptions are vital for progress. We point to the paper by Beringer et al. [173] as an exemplar for how to clearly delineate between what is trusted/untrusted. At the same time, critics should understand that trust assumptions are often necessary to make progress at all.

Our second recommendation concerns the use of metrics. Metrics are useful for tracking progress over time when used appropriately. The HACL* [5] study uses metrics effectively: To quantify verification effort, the authors report proof-to-code ratios and person efforts for various primitives. While these are crude proxies, because the comparison is vertical (same tool, same developers), the numbers sensibly demonstrate that, e.g., code involving bignums requires more work to verify in F^* . Despite their limitations, we argue that even crude metrics (when used appropriately) are better than none for advancing the field. When used inappropriately, however, metrics become dangerous and misleading. Horizontal comparisons across disparate tools tend to be problematic and must be done with care if they are to be used. For example, lines of proof or

analysis times across disparate tools are often incomparable, since modeling a problem in the exact same way is non-trivial.

B. Recommendations to Tool Developers

Although we are still in the early days of seeing verified cryptography deployed in the wild, one major pending challenge is how to make computer-aided cryptography artifacts maintainable. Because computer-aided cryptography tools sit at the bleeding-edge of how cryptography is done, they are constantly evolving, often in non-backwards-compatible ways. When this happens, we must either allow the artifacts (e.g., machine-checked proofs) to become stale, or else muster significant human effort to keep them up to date. Moreover, because cryptography is a moving target, we should expect that even verified implementations (and their proofs) will require updates. This could be to add functionality, or in the worst case, to swiftly patch new vulnerabilities beyond what was verifiably accounted for. To this end, we hope to see more interplay between proof engineering research [184], [185] and computer-aided cryptography research in the coming years.

C. Recommendations to Standardization Bodies

Given its benefits in the TLS 1.3 standardization effort, we believe computer-aided cryptography should play an important role in standardization processes [186]. Traditionally, cryptographic standards are written in a combination of prose, formulas, and pseudocode, and can change drastically between drafts. On top of getting the cryptography right in the first place, standards must also focus on clarity, ease of implementation, and interoperability. Unsurprisingly, standardization processes can be long and arduous. And even when they are successful, the substantial gap between standards and implementations leaves plenty of room for error.

Security proofs can also become a double-edged sword in standardization processes. Proposals supported by hand-written security arguments often cannot be reasonably audited. A plausible claim with a proof that cannot be audited should not be taken as higher assurance than simply stating the claim—we believe the latter is a lesser evil, as it does not create a false sense of security. For example, Hales [187] discusses ill-intentioned security arguments in the context of the Dual EC pseudo-random generator [188]. Another example is the recent discovery of attacks against the AES-OCB2 ISO standard, which was previously believed to be secure [189].

To address these challenges, we advocate the use of computer-aided cryptography, not only to formally certify compliance to standards, but also to facilitate the role of auditors and evaluators in standardization processes, allowing the discussion to focus on the security claims, rather than on whether the supporting security arguments are convincing. We see the current NIST post-quantum standardization effort [190] as an excellent opportunity to put our recommendations into practice, and we encourage the computer-aided cryptography community to engage in the process.

ACKNOWLEDGMENTS

We thank the anonymous reviewers for their useful suggestions; Jason Gross, Boris Köpf, Stever Kremer, Peter Schwabe, and Alwen Tiu for feedback on earlier drafts of the paper; and Tiago Oliveira for help setting up Jasmin and benchmarks.

Work by Manuel Barbosa was supported by National Funds through the Portuguese Foundation for Science and Technology (FCT) under project PTDC/CCI-INF/31698/2017. Work by Gilles Barthe was supported by the Office of Naval Research (ONR) under project N00014-15-1-2750. Work by Karthik Bhargavan was supported by the European Research Council (ERC) under the European Union's Horizon 2020 research and innovation programme (grant agreement no. 683032 - CIRCUS). Work by Bruno Blanchet was supported by the French National Research Agency (ANR) under project TECAP (decision no. ANR-17-CE39-0004-03). Work by Kevin Liao was supported by the National Science Foundation (NSF) through a Graduate Research Fellowship. Work by Bryan Parno was supported by a gift from Bosch, a fellowship from the Alfred P. Sloan Foundation, the NSF under Grant No. 1801369, and the Department of the Navy, Office of Naval Research under Grant No. N00014-18-1-2892.

REFERENCES

- [1] A. Delignat-Lavaud, C. Fournet, M. Kohlweiss, J. Protzenko, A. Rastogi, N. Swamy, S. Z. Béguelin, K. Bhargavan, J. Pan, and J. K. Zinzindohoue, "Implementing and proving the TLS 1.3 record layer," in *IEEE Symposium on Security and Privacy (S&P)*. IEEE Computer Society, 2017, pp. 463–482.
- [2] C. Cremers, M. Horvat, S. Scott, and T. van der Merwe, "Automated analysis and verification of TLS 1.3: 0-rtt, resumption and delayed authentication," in *IEEE Symposium on Security and Privacy (S&P)*. IEEE Computer Society, 2016, pp. 470–485.
- [3] K. Bhargavan, B. Blanchet, and N. Kobeissi, "Verified models and reference implementations for the TLS 1.3 standard candidate," in *IEEE Symposium on Security and Privacy (S&P)*. IEEE Computer Society, 2017, pp. 483–502.
- [4] C. Cremers, M. Horvat, J. Hoyland, S. Scott, and T. van der Merwe, "A comprehensive symbolic analysis of TLS 1.3," in *ACM Conference on Computer and Communications Security (CCS)*. ACM, 2017, pp. 1773–1788.
- [5] J. K. Zinzindohoue, K. Bhargavan, J. Protzenko, and B. Beurdouche, "HACL*: A verified modern cryptographic library," in *ACM Conference on Computer and Communications Security (CCS)*. ACM, 2017, pp. 1789–1806.
- [6] A. Erbsen, J. Philipoom, J. Gross, R. Sloan, and A. Chlipala, "Simple high-level code for cryptographic arithmetic - with proofs, without compromises," in *IEEE Symposium on Security and Privacy (S&P)*. IEEE, 2019, pp. 1202–1219.
- [7] J. Protzenko, B. Parno, A. Fromherz, C. Hawblitzel, M. Polubelova, K. Bhargavan, B. Beurdouche, J. Choi, A. Delignat-Lavaud, C. Fournet, N. Kulatova, T. Ramananandro, A. Rastogi, N. Swamy, C. Wintersteiger, and S. Zanella-Béguelin, "EverCrypt: A fast, verified, cross-platform cryptographic provider," in *IEEE Symposium on Security and Privacy (S&P)*. IEEE, 2020.
- [8] M. Bellare and P. Rogaway, "The security of triple encryption and a framework for code-based game-playing proofs," in *Annual International Conference on the Theory and Applications of Cryptographic Techniques (EUROCRYPT)*, ser. LNCS, vol. 4004. Springer, 2006, pp. 409–426.
- [9] R. Canetti, "Universally composable security: A new paradigm for cryptographic protocols," in *IEEE Annual Symposium on Foundations of Computer Science (FOCS)*. IEEE Computer Society, 2001, pp. 136–145.
- [10] S. Halevi, "A plausible approach to computer-aided cryptographic proofs," *IACR Cryptology ePrint Archive*, vol. 2005, p. 181, 2005.
- [11] K. G. Paterson and G. J. Watson, "Plaintext-dependent decryption: A formal security treatment of SSH-CTR," in *Annual International Conference on the Theory and Applications of Cryptographic Techniques (EUROCRYPT)*, ser. LNCS, vol. 6110. Springer, 2010, pp. 345–361.
- [12] A. Boldyreva, J. P. Degabriele, K. G. Paterson, and M. Stam, "Security of symmetric encryption in the presence of ciphertext fragmentation," in *Annual International Conference on the Theory and Applications of Cryptographic Techniques (EUROCRYPT)*, ser. LNCS, vol. 7237. Springer, 2012, pp. 682–699.
- [13] J. P. Degabriele, K. G. Paterson, and G. J. Watson, "Provable security in the real world," *IEEE Security & Privacy*, vol. 9, no. 3, pp. 33–41, 2011.
- [14] V. Shoup, "Sequences of games: a tool for taming complexity in security proofs," *IACR Cryptology ePrint Archive*, vol. 2004, p. 332, 2004. [Online]. Available: <http://eprint.iacr.org/2004/332>
- [15] Y. Lindell, "How to simulate it - A tutorial on the simulation proof technique," in *Tutorials on the Foundations of Cryptography*. Springer International Publishing, 2017, pp. 277–346.
- [16] S. F. Doghmi, J. D. Guttman, and F. J. Thayer, "Searching for shapes in cryptographic protocols," in *International Conference on Tools and Algorithms for the Construction and Analysis of Systems (TACAS)*, ser. LNCS, vol. 4424. Springer, 2007, pp. 523–537.
- [17] J. Bengtson, K. Bhargavan, C. Fournet, A. D. Gordon, and S. Maffei, "Refinement types for secure implementations," *ACM Trans. Program. Lang. Syst.*, vol. 33, no. 2, pp. 8:1–8:45, 2011.
- [18] M. Backes, C. Hritcu, and M. Maffei, "Union, intersection and refinement types and reasoning about type disjointness for secure protocol implementations," *J. Comput. Secur.*, vol. 22, no. 2, pp. 301–353, Mar. 2014.
- [19] S. Escobar, C. A. Meadows, and J. Meseguer, "Maude-npa: Cryptographic protocol analysis modulo equational properties," in *Foundations of Security Analysis and Design (FOSAD)*, ser. LNCS, vol. 5705. Springer, 2007, pp. 1–50.
- [20] B. Blanchet, "Modeling and verifying security protocols with the applied pi calculus and ProVerif," *Foundations and Trends in Privacy and Security*, vol. 1, no. 1–2, pp. 1–135, Oct. 2016.
- [21] K. Bhargavan, C. Fournet, A. D. Gordon, and S. Tse, "Verified interoperable implementations of security protocols," *ACM Transactions on Programming Languages and Systems*, vol. 31, no. 1, 2008.
- [22] V. Cheval, V. Cortier, and M. Turuani, "A little more conversation, a little less action, a lot more satisfaction: Global states in proverif," in *IEEE Computer Security Foundations Symposium (CSF)*. IEEE Computer Society, 2018, pp. 344–358.
- [23] D. L. Li and A. Tiu, "Combining proverif and automated theorem provers for security protocol verification," in *International Conference on Automated Deduction (CADE)*, ser. LNCS, vol. 11716. Springer, 2019, pp. 354–365.
- [24] M. Arapinis, E. Ritter, and M. D. Ryan, "Statverif: Verification of stateful processes," in *IEEE Computer Security Foundations Symposium (CSF)*. IEEE Computer Society, 2011, pp. 33–47.
- [25] C. J. F. Cremers, "The scyther tool: Verification, falsification, and analysis of security protocols," in *International Conference on Computer-Aided Verification (CAV)*, ser. LNCS, vol. 5123. Springer, 2008, pp. 414–418.
- [26] S. Meier, C. J. F. Cremers, and D. A. Basin, "Strong invariants for the efficient construction of machine-checked protocol security proofs," in *IEEE Computer Security Foundations Symposium (CSF)*. IEEE Computer Society, 2010, pp. 231–245.
- [27] S. Meier, B. Schmidt, C. Cremers, and D. A. Basin, "The TAMARIN prover for the symbolic analysis of security protocols," in *International Conference on Computer-Aided Verification (CAV)*, ser. LNCS, vol. 8044. Springer, 2013, pp. 696–701.
- [28] S. Kremer and R. Künnemann, "Automated analysis of security protocols with global state," in *IEEE Symposium on Security and Privacy (S&P)*. IEEE Computer Society, 2014, pp. 163–178.
- [29] M. Turuani, "The cl-atse protocol analyser," in *International Conference on Term Rewriting and Applications (RTA)*, ser. LNCS, vol. 4098. Springer, 2006, pp. 277–286.
- [30] D. A. Basin, S. Mödersheim, and L. Viganò, "OFMC: A symbolic model checker for security protocols," *Int. J. Inf. Sec.*, vol. 4, no. 3, pp. 181–208, 2005.
- [31] A. Armando and L. Compagna, "SATMC: A sat-based model checker for security protocols," in *European Conference on Logics in Artificial*

- Intelligence (JELIA)*, ser. LNCS, vol. 3229. Springer, 2004, pp. 730–733.
- [32] R. Chadha, V. Cheval, Ștefan Ciobăcă, and S. Kremer, “Automated verification of equivalence properties of cryptographic protocols,” *ACM Trans. Comput. Log.*, vol. 17, no. 4, pp. 23:1–23:32, 2016.
 - [33] V. Cheval, “APTE: an algorithm for proving trace equivalence,” in *International Conference on Tools and Algorithms for the Construction and Analysis of Systems (TACAS)*, ser. LNCS, vol. 8413. Springer, 2014, pp. 587–592.
 - [34] V. Cheval, S. Kremer, and I. Rakotonirina, “DEEPSEC: deciding equivalence properties in security protocols theory and practice,” in *IEEE Symposium on Security and Privacy (S&P)*. IEEE Computer Society, 2018, pp. 529–546.
 - [35] V. Cortier, A. Dallon, and S. Delaune, “Sat-equiv: An efficient tool for equivalence properties,” in *IEEE Computer Security Foundations Symposium (CSF)*. IEEE Computer Society, 2017, pp. 481–494.
 - [36] A. Tiu and J. E. Dawson, “Automating open bisimulation checking for the spi calculus,” in *IEEE Computer Security Foundations Symposium (CSF)*. IEEE Computer Society, 2010, pp. 307–321.
 - [37] J. K. Millen, “A necessarily parallel attack,” in *Workshop on Formal Methods and Security Protocols*, 1999.
 - [38] N. Durgin, P. Lincoln, J. C. Mitchell, and A. Scedrov, “Multiset rewriting and the complexity of bounded security protocols,” *Journal of Computer Security*, vol. 12, no. 2, pp. 247–311, 2004.
 - [39] S. Delaune and L. Hirschi, “A survey of symbolic methods for establishing equivalence-based properties in cryptographic protocols,” *J. Log. Algebr. Meth. Program.*, vol. 87, pp. 127–144, 2017.
 - [40] J. Dreier, C. Duménil, S. Kremer, and R. Sasse, “Beyond subterm-convergent equational theories in automated verification of stateful protocols,” in *International Conference on Principles of Security and Trust (POST)*. Springer-Verlag, 2017.
 - [41] C. Cremers and D. Jackson, “Prime, order please! revisiting small subgroup and invalid curve attacks on protocols using Diffie-Hellman,” in *IEEE Computer Security Foundations Symposium (CSF)*. IEEE, 2019, pp. 78–93.
 - [42] L. C. Paulson, *Isabelle - A Generic Theorem Prover (with a contribution by T. Nipkow)*, ser. LNCS. Springer, 1994, vol. 828.
 - [43] R. Milner, *Communicating and mobile systems - the Pi-calculus*. Cambridge University Press, 1999.
 - [44] M. Abadi and A. D. Gordon, “A calculus for cryptographic protocols: The spi calculus,” in *ACM Conference on Computer and Communications Security (CCS)*. ACM, 1997, pp. 36–47.
 - [45] M. Abadi and C. Fournet, “Mobile values, new names, and secure communication,” in *Symposium on Principles of Programming Languages (POPL)*. ACM, 2001, pp. 104–115.
 - [46] K. Bhargavan, C. Fournet, R. Corin, and E. Zălinescu, “Verified cryptographic implementations for TLS,” *ACM Trans. Inf. Syst. Secur.*, vol. 15, no. 1, pp. 3:1–3:32, 2012.
 - [47] N. Kobeissi, K. Bhargavan, and B. Blanchet, “Automated verification for secure messaging protocols and their implementations: A symbolic and computational approach,” in *IEEE European Symposium on Security and Privacy (EuroS&P)*. IEEE, 2017, pp. 435–450.
 - [48] N. Kobeissi, G. Nicolas, and K. Bhargavan, “Noise explorer: Fully automated modeling and verification for arbitrary noise protocols,” in *IEEE European Symposium on Security and Privacy (EuroS&P)*. IEEE, 2019, pp. 356–370.
 - [49] D. A. Basin, J. Dreier, L. Hirschi, S. Radomirovic, R. Sasse, and V. Stettler, “A formal analysis of 5g authentication,” in *ACM Conference on Computer and Communications Security (CCS)*. ACM, 2018, pp. 1383–1396.
 - [50] C. Cremers, M. Dehnel-Wild, and K. Milner, “Secure authentication in the grid: A formal analysis of DNP3 SAV5,” *Journal of Computer Security*, vol. 27, no. 2, pp. 203–232, 2019.
 - [51] G. Girol, L. Hirschi, R. Sasse, D. Jackson, C. Cremers, and D. Basin, “A Spectral Analysis of Noise: A Comprehensive, Automated, Formal Analysis of Diffie-Hellman Protocols,” in *Proc. of USENIX Security*, 2020.
 - [52] B. Blanchet, M. Abadi, and C. Fournet, “Automated verification of selected equivalences for security protocols,” *Journal of Logic and Algebraic Programming*, vol. 75, no. 1, pp. 3–51, Feb.–Mar. 2008.
 - [53] S. Santiago, S. Escobar, C. Meadows, and J. Meseguer, “A formal definition of protocol indistinguishability and its verification using Maude-NPA,” in *Security and Trust Management (STM)*, ser. LNCS, vol. 8743. Berlin, Heidelberg: Springer, Sep. 2014, pp. 162–177.
 - [54] D. Basin, J. Dreier, and R. Casse, “Automated symbolic proofs of observational equivalence,” in *ACM Conference on Computer and Communications Security (CCS)*. New York, NY: ACM Press, Oct. 2015, pp. 1144–1155.
 - [55] G. Barthe, B. Grégoire, and B. Schmidt, “Automated proofs of pairing-based cryptography,” in *ACM Conference on Computer and Communications Security (CCS)*. ACM, 2015, pp. 1156–1168.
 - [56] G. Barthe, B. Grégoire, and S. Z. Béguelin, “Formal certification of code-based cryptographic proofs,” in *Symposium on Principles of Programming Languages (POPL)*. ACM, 2009, pp. 90–101.
 - [57] D. A. Basin, A. Lochbihler, and S. R. Sefidgar, “CryptHOL: Game-based proofs in higher-order logic,” *IACR Cryptology ePrint Archive*, vol. 2017, p. 753, 2017.
 - [58] B. Blanchet, “A computationally sound mechanized prover for security protocols,” *IEEE Transactions on Dependable and Secure Computing*, vol. 5, no. 4, pp. 193–207, Oct.–Dec. 2008.
 - [59] G. Barthe, B. Grégoire, S. Heraud, and S. Z. Béguelin, “Computer-aided security proofs for the working cryptographer,” in *International Cryptology Conference (CRYPTO)*, ser. LNCS, vol. 6841. Springer, 2011, pp. 71–90.
 - [60] N. Swamy, C. Hritcu, C. Keller, A. Rastogi, A. Delignat-Lavaud, S. Forest, K. Bhargavan, C. Fournet, P. Strub, M. Kohlweiss, J. K. Zinzindohoue, and S. Z. Béguelin, “Dependent types and monadic effects in F,” in *Symposium on Principles of Programming Languages (POPL)*. ACM, 2016, pp. 256–270.
 - [61] A. Petcher and G. Morrisett, “The foundational cryptography framework,” in *International Conference on Principles of Security and Trust (POST)*, ser. LNCS, vol. 9036. Springer, 2015, pp. 53–72.
 - [62] G. Barthe, J. M. Crespo, B. Grégoire, C. Kunz, Y. Lakhnech, B. Schmidt, and S. Z. Béguelin, “Fully automated analysis of padding-based encryption in the computational model,” in *ACM Conference on Computer and Communications Security (CCS)*. ACM, 2013, pp. 1247–1260.
 - [63] “The coq proof assistant.” [Online]. Available: <https://coq.inria.fr/>
 - [64] “Isabelle.” [Online]. Available: <https://isabelle.in.tum.de/>
 - [65] C. W. Barrett and C. Tinelli, “Satisfiability modulo theories,” in *Handbook of Model Checking*. Springer, 2018, pp. 305–343.
 - [66] L. M. de Moura and N. Björner, “Z3: an efficient SMT solver,” in *International Conference on Tools and Algorithms for the Construction and Analysis of Systems (TACAS)*, ser. LNCS, vol. 4963. Springer, 2008, pp. 337–340.
 - [67] B. Lipp, B. Blanchet, and K. Bhargavan, “A mechanised cryptographic proof of the wireguard virtual private network protocol,” in *IEEE European Symposium on Security and Privacy (EuroS&P)*. IEEE, 2019, pp. 231–246.
 - [68] J. B. Almeida, M. Barbosa, G. Barthe, M. Campagna, E. Cohen, B. Grégoire, V. Pereira, B. Portela, P. Strub, and S. Tasiran, “A machine-checked proof of security for AWS key management service,” in *ACM Conference on Computer and Communications Security (CCS)*. ACM, 2019, pp. 63–78.
 - [69] J. B. Almeida, C. Baritel-Ruet, M. Barbosa, G. Barthe, F. Dupressoir, B. Grégoire, V. Laporte, T. Oliveira, A. Stoughton, and P. Strub, “Machine-checked proofs for cryptographic standards: Indifferentiability of sponge and secure high-assurance implementations of SHA-3,” in *ACM Conference on Computer and Communications Security (CCS)*. ACM, 2019, pp. 1607–1622.
 - [70] K. Bhargavan, C. Fournet, M. Kohlweiss, A. Pironti, and P. Strub, “Implementing TLS with verified cryptographic security,” in *IEEE Symposium on Security and Privacy (S&P)*. IEEE Computer Society, 2013, pp. 445–459.
 - [71] K. Bhargavan, C. Fournet, M. Kohlweiss, A. Pironti, P.-Y. Strub, and S. Zanella-Béguelin, “Proving the TLS handshake secure (as it is),” in *International Cryptology Conference (CRYPTO)*, 2014.
 - [72] K. Bhargavan, A. Delignat-Lavaud, C. Fournet, A. Pironti, and P. Strub, “Triple handshakes and cookie cutters: Breaking and fixing authentication over TLS,” in *IEEE Symposium on Security and Privacy (S&P)*. IEEE Computer Society, 2014, pp. 98–113.
 - [73] B. Beurdouche, K. Bhargavan, A. Delignat-Lavaud, C. Fournet, M. Kohlweiss, A. Pironti, P. Strub, and J. K. Zinzindohoue, “A messy state of the union: Taming the composite state machines of TLS,” in *IEEE Symposium on Security and Privacy (S&P)*. IEEE Computer Society, 2015, pp. 535–552.
 - [74] C. E. Landwehr, D. Boneh, J. C. Mitchell, S. M. Bellovin, S. Landau,

- and M. E. Lesk, "Privacy and cybersecurity: The next 100 years," *Proc. of the IEEE*, vol. 100, no. Centennial-Issue, pp. 1659–1673, 2012.
- [75] K. Liao, M. A. Hammer, and A. Miller, "ILC: a calculus for composable, computational cryptography," in *ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI)*. ACM, 2019, pp. 640–654.
 - [76] R. Canetti, A. Stoughton, and M. Varia, "EasyUC: Using EasyCrypt to mechanize proofs of universally composable security," in *IEEE Computer Security Foundations Symposium (CSF)*. IEEE, 2019, pp. 167–183.
 - [77] A. Lochbihler, S. R. Sefidgar, D. A. Basin, and U. Maurer, "Formalizing constructive cryptography using CryptHOL," in *IEEE Computer Security Foundations Symposium (CSF)*. IEEE, 2019, pp. 152–166.
 - [78] J. A. Akinyele, M. Green, and S. Hohenberger, "Using SMT solvers to automate design tasks for encryption and signature schemes," in *2013 ACM SIGSAC Conference on Computer and Communications Security, CCS'13, Berlin, Germany, November 4-8, 2013*. ACM, 2013, pp. 399–410.
 - [79] A. J. Malozemoff, J. Katz, and M. D. Green, "Automated analysis and synthesis of block-cipher modes of operation," in *IEEE Computer Security Foundations Symposium (CSF)*. IEEE Computer Society, 2014, pp. 140–152.
 - [80] V. T. Hoang, J. Katz, and A. J. Malozemoff, "Automated analysis and synthesis of authenticated encryption schemes," in *ACM Conference on Computer and Communications Security (CCS)*. ACM, 2015, pp. 84–95.
 - [81] G. Barthe, E. Fagerholm, D. Fiore, A. Scedrov, B. Schmidt, and M. Tibouchi, "Strongly-optimal structure preserving signatures from type II pairings: synthesis and lower bounds," *IET Information Security*, vol. 10, no. 6, pp. 358–371, 2016.
 - [82] B. Parno, J. Howell, C. Gentry, and M. Raykova, "Pinocchio: nearly practical verifiable computation," *Commun. ACM*, vol. 59, no. 2, pp. 103–112, 2016.
 - [83] C. Costello, C. Fournet, J. Howell, M. Kohlweiss, B. Kreuter, M. Naehrig, B. Parno, and S. Zahur, "Geppetto: Versatile verifiable computation," in *IEEE Symposium on Security and Privacy (S&P)*, 2015, pp. 253–270.
 - [84] S. T. V. Setty, V. Vu, N. Panpalia, B. Braun, A. J. Blumberg, and M. Walfish, "Taking proof-based verified computation a few steps closer to practicality," in *USENIX Security Symposium (USENIX)*. USENIX Association, 2012, pp. 253–268.
 - [85] E. Ben-Sasson, A. Chiesa, D. Genkin, E. Tromer, and M. Virza, "SNARKs for C: verifying program executions succinctly and in zero knowledge," in *International Cryptology Conference (CRYPTO)*, ser. LNCS, vol. 8043. Springer, 2013, pp. 90–108.
 - [86] J. B. Almeida, E. Bangerter, M. Barbosa, S. Krenn, A. Sadeghi, and T. Schneider, "A certifying compiler for zero-knowledge proofs of knowledge based on sigma-protocols," in *European Symposium on Research in Computer Security (ESORICS)*, 2010, pp. 151–167.
 - [87] M. Fredrikson and B. Livshits, "Z0: An optimizing distributing zero-knowledge compiler," in *USENIX Security Symposium (USENIX)*, 2014, pp. 909–924.
 - [88] S. Meiklejohn, C. C. Erway, A. Küpçü, T. Hinkle, and A. Lysyanskaya, "ZKPD: A language-based system for efficient zero-knowledge proofs and electronic cash," in *USENIX Security Symposium (USENIX)*. USENIX Association, 2010, pp. 193–206.
 - [89] M. Backes, M. Maffei, and K. Pecina, "Automated synthesis of secure distributed applications," in *Symposium on Network and Distributed System Security (NDSS)*. The Internet Society, 2012.
 - [90] M. Hastings, B. Hemenway, D. Noble, and S. Zdancewic, "Sok: General purpose compilers for secure multi-party computation," in *IEEE Symposium on Security and Privacy (S&P)*, 2019, pp. 1220–1237.
 - [91] J. B. Almeida, M. Barbosa, E. Bangerter, G. Barthe, S. Krenn, and S. Z. Béguélin, "Full proof cryptography: verifiable compilation of efficient zero-knowledge protocols," in *ACM Conference on Computer and Communications Security (CCS)*. ACM, 2012, pp. 488–500.
 - [92] J. B. Almeida, M. Barbosa, G. Barthe, F. Dupressoir, B. Grégoire, V. Laporte, and V. Pereira, "A fast and verified software stack for secure function evaluation," in *ACM Conference on Computer and Communications Security (CCS)*. ACM, 2017, pp. 1989–2006.
 - [93] C. Fournet, C. Keller, and V. Laporte, "A certified compiler for verifiable computing," in *IEEE Computer Security Foundations Symposium (CSF)*, 2016, pp. 268–280.
 - [94] A. Rastogi, N. Swamy, and M. Hicks, "Wys*: A DSL for verified secure multi-party computations," in *International Conference on Principles of Security and Trust (POST)*, 2019, pp. 99–122.
 - [95] B. Blanchet, "Security protocol verification: Symbolic and computational models," in *International Conference on Principles of Security and Trust (POST)*, ser. LNCS, vol. 7215. Springer, 2012, pp. 3–29.
 - [96] V. Cortier, S. Kremer, and B. Warinschi, "A survey of symbolic methods in computational analysis of cryptographic systems," *J. Autom. Reasoning*, vol. 46, no. 3-4, pp. 225–259, 2011.
 - [97] R. Dockins, A. Foltzer, J. Hendrix, B. Huffman, D. McNamee, and A. Tomb, "Constructing semantic models of programs with the software analysis workbench," in *International Conference on Verified Software Theories, Tools, and Experiments (VSTTE)*, ser. LNCS, vol. 9971, 2016, pp. 56–72.
 - [98] Y. Fu, J. Liu, X. Shi, M. Tsai, B. Wang, and B. Yang, "Signed cryptographic program verification with typed cryptoline," in *ACM Conference on Computer and Communications Security (CCS)*. ACM, 2019, pp. 1591–1606.
 - [99] K. R. M. Leino, "Dafny: An automatic program verifier for functional correctness," in *International Conference on Logic for Programming, Artificial Intelligence, and Reasoning (LPAR)*, ser. LNCS, vol. 6355. Springer, 2010, pp. 348–370.
 - [100] P. Cuoq, F. Kirchner, N. Kosmatov, V. Prevosto, J. Signoles, and B. Yakobowski, "Frama-c - A software analysis perspective," in *International Conference on Software Engineering and Formal Methods (SEFM)*, ser. LNCS, vol. 7504. Springer, 2012, pp. 233–247.
 - [101] D. J. Bernstein and P. Schwabe, "gfverif: Fast and easy verification of finite-field arithmetic," 2016. [Online]. Available: <http://gfverif.cryptojedi.org>
 - [102] J. B. Almeida, M. Barbosa, G. Barthe, A. Blot, B. Grégoire, V. Laporte, T. Oliveira, H. Pacheco, B. Schmidt, and P. Strub, "Jasmin: High-assurance and high-speed cryptography," in *ACM Conference on Computer and Communications Security (CCS)*. ACM, 2017, pp. 1807–1823.
 - [103] B. Bond, C. Hawblitzel, M. Kapritsos, K. R. M. Leino, J. R. Lorch, B. Parno, A. Rane, S. T. V. Setty, and L. Thompson, "Vale: Verifying high-performance cryptographic assembly code," in *USENIX Security Symposium (USENIX)*. USENIX Association, 2017, pp. 917–934.
 - [104] A. Fromherz, N. Giannarakis, C. Hawblitzel, B. Parno, A. Rastogi, and N. Swamy, "A verified, efficient embedding of a verifiable assembly language," *PACMPL*, vol. 3, no. POPL, pp. 63:1–63:30, 2019.
 - [105] A. W. Appel, "Verified software toolchain - (invited talk)," in *European Symposium on Programming (ESOP)*, ser. LNCS, vol. 6602. Springer, 2011, pp. 1–17.
 - [106] J. Filiâtre and A. Paskevich, "Why3 - where programs meet provers," in *European Symposium on Programming (ESOP)*, ser. LNCS, vol. 7792. Springer, 2013, pp. 125–128.
 - [107] D. J. Bernstein and T. Lange, "ebacs: Ecrypt benchmarking of cryptographic systems," 2009. [Online]. Available: <https://bench.cr.yp.to>
 - [108] Z. Durumeric, J. Kasten, D. Adrian, J. A. Halderman, M. Bailey, F. Li, N. Weaver, J. Amann, J. Beekman, M. Payer, and V. Paxson, "The matter of heartbleed," in *Internet Measurement Conference (IMC)*. ACM, 2014, pp. 475–488.
 - [109] S. Gueron and V. Krasnov, "The fragility of AES-GCM authentication algorithm," in *Proc. of the Conference on Information Technology: New Generations*, Apr. 2014.
 - [110] B. B. Brumley, M. Barbosa, D. Page, and F. Vercauteren, "Practical realisation and elimination of an ecc-related software bug attack," in *Cryptographers' Track at the RSA Conference (CT-RSA)*, ser. LNCS, vol. 7178. Springer, 2012, pp. 171–186.
 - [111] X. Leroy, "Formal verification of a realistic compiler," *Commun. ACM*, vol. 52, no. 7, pp. 107–115, 2009.
 - [112] T. Oliveira, J. L. Hernandez, H. Hisil, A. Faz-Hernández, and F. Rodríguez-Henríquez, "How to (pre-)compute a ladder - improving the performance of X25519 and X448," in *International Conference on Selected Areas in Cryptography (SAC)*, ser. LNCS, vol. 10719. Springer, 2017, pp. 172–191.
 - [113] T. Chou, "Sandy2x: New curve25519 speed records," in *International Conference on Selected Areas in Cryptography (SAC)*, ser. LNCS, vol. 9566. Springer, 2015, pp. 145–160.
 - [114] Y. Chen, C. Hsu, H. Lin, P. Schwabe, M. Tsai, B. Wang, B. Yang, and S. Yang, "Verifying curve25519 software," in *ACM Conference*

- on *Computer and Communications Security (CCS)*. ACM, 2014, pp. 299–309.
- [115] “curve25519-donna: Implementations of a fast elliptic-curve Diffie-Hellman primitive,” <https://github.com/agl/curve25519-donna>.
- [116] D. J. Bernstein, “Curve25519: New Diffie-Hellman speed records,” in *IACR International Conference on Practice and Theory of Public-Key Cryptography (PKC)*, ser. LNCS, vol. 3958. Springer, 2006, pp. 207–228.
- [117] J. B. Almeida, M. Barbosa, G. Barthe, B. Grégoire, A. Koutsos, V. Laporte, T. Oliveira, and P. Strub, “The last mile: High-assurance and high-speed cryptographic implementations,” *CoRR*, vol. abs/1904.04606, 2019.
- [118] J. P. Lim and S. Nagarakatte, “Automatic equivalence checking for assembly implementations of cryptography libraries,” in *Proc. of the IEEE/ACM International Symposium on Code Generation and Optimization, (CGO)*. IEEE, 2019, pp. 37–49.
- [119] P. L. Montgomery, “Modular multiplication without trial division,” *Mathematics of computation*, vol. 44, no. 170, pp. 519–521, 1985.
- [120] “The GNU Multiple Precision Arithmetic Library.” [Online]. Available: <https://gmplib.org/>
- [121] G. Klein, J. Andronick, K. Elphinstone, T. C. Murray, T. Sewell, R. Kolanski, and G. Heiser, “Comprehensive formal verification of an OS microkernel,” *ACM Trans. Comput. Syst.*, vol. 32, no. 1, pp. 2:1–2:70, 2014.
- [122] R. Gu, J. Koenig, T. Ramanandro, Z. Shao, X. N. Wu, S. Weng, H. Zhang, and Y. Guo, “Deep specifications and certified abstraction layers,” in *Symposium on Principles of Programming Languages (POPL)*. ACM, 2015, pp. 595–608.
- [123] H. Mai, E. Pek, H. Xue, S. T. King, and P. Madhusudan, “Verifying security invariants in ExpressOS,” in *International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS)*. ACM, 2013, pp. 293–304.
- [124] G. Morrisett, G. Tan, J. Tassarotti, J. Tristan, and E. Gan, “Rocksalt: better, faster, stronger SFI for the x86,” in *ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI)*. ACM, 2012, pp. 395–404.
- [125] X. Wang, D. Lazar, N. Zeldovich, A. Chlipala, and Z. Tatlock, “Jitk: A trustworthy in-kernel interpreter infrastructure,” in *USENIX Conference on Operating Systems Design and Implementation (OSDI)*. USENIX Association, 2014, pp. 33–47.
- [126] H. Chen, D. Ziegler, T. Chajed, A. Chlipala, M. F. Kaashoek, and N. Zeldovich, “Using crash hoare logic for certifying the FSCQ file system,” in *ACM Symposium on Operating Systems Principles (SOSP)*. ACM, 2015, pp. 18–37.
- [127] A. Vasudevan, S. Chaki, L. Jia, J. M. McCune, J. Newsome, and A. Datta, “Design, implementation and verification of an extensible and modular hypervisor framework,” in *IEEE Symposium on Security and Privacy (S&P)*. IEEE Computer Society, 2013, pp. 430–444.
- [128] J. R. Wilcox, D. Woos, P. Panthekha, Z. Tatlock, X. Wang, M. D. Ernst, and T. E. Anderson, “Verdi: a framework for implementing and formally verifying distributed systems,” in *ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI)*. ACM, 2015, pp. 357–368.
- [129] O. Padon, K. L. McMillan, A. Panda, M. Sagiv, and S. Shoham, “Ivy: safety verification by interactive generalization,” in *ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI)*. ACM, 2016, pp. 614–630.
- [130] C. Hawblitzel, J. Howell, M. Kapritsos, J. R. Lorch, B. Parno, M. L. Roberts, S. T. V. Setty, and B. Zill, “Ironfleet: proving practical distributed systems correct,” in *ACM Symposium on Operating Systems Principles (SOSP)*. ACM, 2015, pp. 1–17.
- [131] C. Hawblitzel, J. Howell, J. R. Lorch, A. Narayan, B. Parno, D. Zhang, and B. Zill, “Ironclad apps: End-to-end security via automated full-system verification,” in *USENIX Conference on Operating Systems Design and Implementation (OSDI)*. USENIX Association, 2014, pp. 165–181.
- [132] J. B. Almeida, M. Barbosa, J. S. Pinto, and B. Vieira, “Formal verification of side-channel countermeasures using self-composition,” *Sci. Comput. Program.*, vol. 78, no. 7, pp. 796–812, 2013.
- [133] G. Doychev, D. Feld, B. Köpf, L. Mauborgne, and J. Reineke, “Cacheaudit: A tool for the static analysis of cache side channels,” in *USENIX Security Symposium (USENIX)*. USENIX Association, 2013, pp. 431–446.
- [134] J. B. Almeida, M. Barbosa, G. Barthe, F. Dupressoir, and M. Emmi, “Verifying constant-time implementations,” in *USENIX Security Symposium (USENIX)*. USENIX Association, 2016, pp. 53–70.
- [135] C. Watt, J. Renner, N. Popescu, S. Cauligi, and D. Stefan, “Ct-wasm: type-driven secure cryptography for the web ecosystem,” *PACMPL*, vol. 3, no. POPL, pp. 77:1–77:29, 2019.
- [136] S. Cauligi, G. Soeller, B. Johannesmeyer, F. Brown, R. S. Wahby, J. Renner, B. Grégoire, G. Barthe, R. Jhala, and D. Stefan, “Fact: a DSL for timing-sensitive computation,” in *ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI)*. ACM, 2019, pp. 174–189.
- [137] B. Rodrigues, F. M. Q. Pereira, and D. F. Aranha, “Sparse representation of implicit flows with applications to side-channel detection,” in *International Conference on Compiler Construction (CC)*. ACM, 2016, pp. 110–120.
- [138] B. Köpf, L. Mauborgne, and M. Ochoa, “Automatic quantification of cache side-channels,” in *International Conference on Computer-Aided Verification (CAV)*, ser. LNCS, vol. 7358. Springer, 2012, pp. 564–580.
- [139] J. Protzenko, J. K. Zinzindohoué, A. Rastogi, T. Ramanandro, P. Wang, S. Z. Béguelin, A. Delignat-Lavaud, C. Hritcu, K. Bhargavan, C. Fournet, and N. Swamy, “Verified low-level programming embedded in F,” *PACMPL*, vol. 1, no. ICFP, pp. 17:1–17:29, 2017.
- [140] M. Wu, S. Guo, P. Schaumont, and C. Wang, “Eliminating timing side-channel leaks using program repair,” in *International Symposium on Software Testing and Analysis (ISSTA)*. ACM, 2018, pp. 15–26.
- [141] G. Barthe, G. Betarte, J. D. Campo, C. D. Luna, and D. Pichardie, “System-level non-interference for constant-time cryptography,” in *ACM Conference on Computer and Communications Security (CCS)*. ACM, 2014, pp. 1267–1279.
- [142] D. Brumley and D. Boneh, “Remote timing attacks are practical,” in *USENIX Security Symposium (USENIX)*. USENIX Association, 2003.
- [143] D. J. Bernstein, “Cache-timing attacks on AES,” 2005.
- [144] J.-P. Aumasson, “Guidelines for Low-Level Cryptography Software,” <https://github.com/veorq/cryptocoding>.
- [145] B. Moller, “Security of CBC ciphersuites in SSL/TLS: Problems and countermeasures,” 2004. [Online]. Available: <http://www.openssl.org/~bodo/tls-cbc.txt>
- [146] N. J. AlFardan and K. G. Paterson, “Lucky thirteen: Breaking the TLS and DTLS record protocols,” in *IEEE Symposium on Security and Privacy (S&P)*. IEEE Computer Society, 2013, pp. 526–540.
- [147] J. Somorovsky, “Curious padding oracle in OpenSSL (cve-2016-2107),” 2016. [Online]. Available: <https://web-in-security.blogspot.com/2016/05/curious-padding-oracle-in-openssl-cve.html>
- [148] P. Kocher, J. Horn, A. Fogh, D. Genkin, D. Gruss, W. Haas, M. Hamburg, M. Lipp, S. Mangard, T. Prescher, M. Schwarz, and Y. Yarom, “Spectre attacks: Exploiting speculative execution,” in *IEEE Symposium on Security and Privacy (S&P)*. IEEE, 2019, pp. 1–19.
- [149] M. Lipp, M. Schwarz, D. Gruss, T. Prescher, W. Haas, A. Fogh, J. Horn, S. Mangard, P. Kocher, D. Genkin, Y. Yarom, and M. Hamburg, “Meltdown: Reading kernel memory from user space,” in *USENIX Security Symposium (USENIX)*. USENIX Association, 2018, pp. 973–990.
- [150] G. Barthe, B. Grégoire, and V. Laporte, “Secure compilation of side-channel countermeasures: The case of cryptographic “constant-time,”” in *IEEE Computer Security Foundations Symposium (CSF)*. IEEE Computer Society, 2018, pp. 328–343.
- [151] D. Molnar, M. Piotrowski, D. Schultz, and D. A. Wagner, “The program counter security model: Automatic detection and removal of control-flow side channel attacks,” in *International Conference on Information Security and Cryptology (ICISC)*, ser. LNCS, vol. 3935. Springer, 2005, pp. 156–168.
- [152] A. Langley, “ctgrind,” 2010. [Online]. Available: <https://github.com/agl/ctgrind/>
- [153] M. Andryscio, A. Nötzli, F. Brown, R. Jhala, and D. Stefan, “Towards verified, constant-time floating point operations,” in *ACM Conference on Computer and Communications Security (CCS)*. ACM, 2018, pp. 1369–1382.
- [154] M. Andryscio, D. Kohlbrenner, K. Mowery, R. Jhala, S. Lerner, and H. Shacham, “On subnormal floating point and abnormal timing,” in *IEEE Symposium on Security and Privacy (S&P)*. IEEE Computer Society, 2015, pp. 623–639.
- [155] D. Kohlbrenner and H. Shacham, “On the effectiveness of mitigations against floating-point timing channels,” in *USENIX Security Symposium (USENIX)*. USENIX Association, 2017, pp. 69–81.

- [156] T. Kaufmann, H. Pelletier, S. Vaudenay, and K. Villegas, "When constant-time source yields variable-time binary: Exploiting curve25519-donna built with MSVC 2015," in *International Conference on Cryptology and Network Security (CANS)*, ser. LNCS, vol. 10052, 2016, pp. 573–582.
- [157] G. Barthe, S. Blazy, B. Grégoire, R. Hutin, V. Laporte, D. Pichardie, and A. Trieu, "Formal verification of a constant-time preserving C compiler," *Proc. ACM Program. Lang.*, vol. 4, no. POPL, pp. 7:1–7:30, 2020.
- [158] A. Reid, "Trustworthy specifications of arm® v8-a and v8-m system level architecture," in *2016 Formal Methods in Computer-Aided Design, FMCAD 2016, Mountain View, CA, USA, October 3-6, 2016*. IEEE, 2016, pp. 161–168.
- [159] A. Armstrong, T. Bauereiss, B. Campbell, A. Reid, K. E. Gray, R. M. Norton, P. Mundkur, M. Wassell, J. French, C. Pulte, S. Flur, I. Stark, N. Krishnaswami, and P. Sewell, "ISA semantics for armv8-a, risc-v, and CHERI-MIPS," *PACMPL*, vol. 3, no. POPL, pp. 71:1–71:31, 2019.
- [160] G. Heiser, "For safety's sake: We need a new hardware-software contract!" *IEEE Design & Test*, vol. 35, no. 2, pp. 27–30, 2018.
- [161] D. Zhang, Y. Wang, G. E. Suh, and A. C. Myers, "A hardware design language for timing-sensitive information-flow security," in *International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS)*. ACM, 2015, pp. 503–516.
- [162] M. Tiwari, H. M. G. Wassel, B. Mazloom, S. Mysore, F. T. Chong, and T. Sherwood, "Complete information flow tracking from the gates up," in *International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS)*. ACM, 2009, pp. 109–120.
- [163] X. Li, V. Kashyap, J. K. Oberg, M. Tiwari, V. R. Rajarathinam, R. Kastner, T. Sherwood, B. Hardekopf, and F. T. Chong, "Sapper: a language for hardware-level security policy enforcement," in *International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS)*. ACM, 2014, pp. 97–112.
- [164] X. Li, M. Tiwari, J. Oberg, V. Kashyap, F. T. Chong, T. Sherwood, and B. Hardekopf, "Caisson: a hardware description language for secure information flow," in *ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI)*. ACM, 2011, pp. 109–120.
- [165] K. von Gleissenthall, R. G. Kici, D. Stefan, and R. Jhala, "IODINE: verifying constant-time execution of hardware," in *USENIX Security Symposium (USENIX)*. USENIX Association, 2019, pp. 1411–1428.
- [166] H. Eldib, C. Wang, and P. Schaumont, "Smt-based verification of software countermeasures against side-channel attacks," in *International Conference on Tools and Algorithms for the Construction and Analysis of Systems (TACAS)*, ser. LNCS, vol. 8413. Springer, 2014, pp. 62–77.
- [167] A. G. Bayrak, F. Regazzoni, D. Novo, and P. Ienne, "Sleuth: Automated verification of software power analysis countermeasures," in *Conference on Cryptographic Hardware and Embedded Systems (CHES)*, ser. LNCS, vol. 8086. Springer, 2013, pp. 293–310.
- [168] A. Moss, E. Oswald, D. Page, and M. Tunstall, "Compiler assisted masking," in *Conference on Cryptographic Hardware and Embedded Systems (CHES)*, ser. LNCS, vol. 7428. Springer, 2012, pp. 58–75.
- [169] H. Eldib and C. Wang, "Synthesis of masking countermeasures against side channel attacks," in *International Conference on Computer-Aided Verification (CAV)*, ser. LNCS, vol. 8559. Springer, 2014, pp. 114–130.
- [170] G. Barthe, S. Belaïd, F. Dupressoir, P. Fouque, B. Grégoire, and P. Strub, "Verified proofs of higher-order masking," in *Annual International Conference on the Theory and Applications of Cryptographic Techniques (EUROCRYPT)*, ser. LNCS, vol. 9056. Springer, 2015, pp. 457–485.
- [171] G. Barthe, S. Belaïd, G. Cassiers, P. Fouque, B. Grégoire, and F. Standardt, "maskverif: Automated verification of higher-order masking in presence of physical defaults," in *European Symposium on Research in Computer Security (ESORICS)*, ser. LNCS, vol. 11735. Springer, 2019, pp. 300–318.
- [172] J. B. Almeida, M. Barbosa, G. Barthe, and F. Dupressoir, "Certified computer-aided cryptography: efficient provably secure machine code from high-level implementations," in *ACM Conference on Computer and Communications Security (CCS)*. ACM, 2013, pp. 1217–1230.
- [173] L. Beringer, A. Petcher, K. Q. Ye, and A. W. Appel, "Verified correctness and security of openssl HMAC," in *USENIX Security Symposium (USENIX)*. USENIX Association, 2015, pp. 207–221.
- [174] J. B. Almeida, M. Barbosa, G. Barthe, and F. Dupressoir, "Verifiable side-channel security of cryptographic implementations: Constant-time MEE-CBC," in *International Conference on Fast Software Encryption (FSE)*, ser. LNCS, vol. 9783. Springer, 2016, pp. 163–184.
- [175] A. Tomb, "Automated verification of real-world cryptographic implementations," *IEEE Security & Privacy*, vol. 14, no. 6, pp. 26–33, 2016.
- [176] J. K. Zinzindohoue, E. Bartzia, and K. Bhargavan, "A verified extensible library of elliptic curves," in *IEEE Computer Security Foundations Symposium (CSF)*. IEEE Computer Society, 2016, pp. 296–309.
- [177] K. Q. Ye, M. Green, N. Sanguansin, L. Beringer, A. Petcher, and A. W. Appel, "Verified correctness and security of mbedTLS HMAC-DRBG," in *ACM Conference on Computer and Communications Security (CCS)*. ACM, 2017, pp. 2007–2020.
- [178] A. Chudnov, N. Collins, B. Cook, J. Dodds, B. Huffman, C. MacCárthaigh, S. Magill, E. Mertens, E. Mullen, S. Tasiran, A. Tomb, and E. Westbrook, "Continuous formal verification of amazon s2n," in *International Conference on Computer-Aided Verification (CAV)*, ser. LNCS, vol. 10982. Springer, 2018, pp. 430–446.
- [179] K. Eldefrawy and V. Pereira, "A high-assurance evaluator for machine-checked secure multiparty computation," in *ACM Conference on Computer and Communications Security (CCS)*. ACM, 2019, pp. 851–868.
- [180] J. Protzenko, B. Beurdouche, D. Merigoux, and K. Bhargavan, "Formally verified cryptographic web applications in webassembly," in *IEEE Symposium on Security and Privacy (S&P)*. IEEE, 2019, pp. 1256–1274.
- [181] C. Meyer and J. Schwenk, "Sok: Lessons learned from SSL/TLS attacks," in *Proc. of the International Workshop on Information Security Applications (WISA)*, ser. LNCS, vol. 8267. Springer, 2013, pp. 189–209.
- [182] J. Clark and P. C. van Oorschot, "Sok: SSL and HTTPS: revisiting past challenges and evaluating certificate trust model enhancements," in *IEEE Symposium on Security and Privacy (S&P)*. IEEE Computer Society, 2013, pp. 511–525.
- [183] K. G. Paterson and T. van der Merwe, "Reactive and proactive standardisation of TLS," in *International Conference on Security Standardisation Research (SSR)*, ser. LNCS, vol. 10074. Springer, 2016, pp. 160–186.
- [184] T. Ringer, K. Palmskog, I. Sergey, M. Gligoric, and Z. Tatlock, "QED at large: A survey of engineering of formally verified software," *Foundations and Trends in Programming Languages*, vol. 5, no. 2-3, pp. 102–281, 2019.
- [185] D. R. Jeffery, M. Staples, J. Andronick, G. Klein, and T. C. Murray, "An empirical research agenda for understanding formal methods productivity," *Information & Software Technology*, vol. 60, pp. 102–112, 2015.
- [186] K. Bhargavan, F. Kiefer, and P. Strub, "hacspect: Towards verifiable crypto standards," in *International Conference on Security Standardisation Research (SSR)*, ser. LNCS, vol. 11322. Springer, 2018, pp. 1–20.
- [187] T. C. Hales, "The nsa back door to nist," *Notices of the AMS*, vol. 61, no. 2, pp. 190–192, 2014.
- [188] S. Checkoway, J. Maskiewicz, C. Garman, J. Fried, S. Cohny, M. Green, N. Heninger, R. Weinmann, E. Rescorla, and H. Shacham, "A systematic analysis of the juniper dual EC incident," in *ACM Conference on Computer and Communications Security (CCS)*. ACM, 2016, pp. 468–479.
- [189] A. Inoue, T. Iwata, K. Minematsu, and B. Poettering, "Cryptanalysis of OCB2: attacks on authenticity and confidentiality," in *International Cryptology Conference (CRYPTO)*, 2019, pp. 3–31.
- [190] L. Chen, L. Chen, S. Jordan, Y.-K. Liu, D. Moody, R. Peralta, R. Perlner, and D. Smith-Tone, *Report on post-quantum cryptography*. US Department of Commerce, National Institute of Standards and Technology, 2016.

Provable Security in the Real World

Provable security plays an important role in the design and analysis of systems using cryptography. However, protocols can be vulnerable to attacks outside the scope of the existing formal analyses.

If we define *science* as an objective approach to analysis and pursuit of knowledge on the basis of rigorous logical arguments, then cryptography undoubtedly has become a science. The area of research responsible for bringing about this transition is commonly known as *provable security*. Provable security introduces formal definitions of security and adopts techniques from probability theory and computational complexity theory to analyze the security of cryptographic constructs. Recently, the merits of this approach have been subject to debate, spurred in part by a series of articles by Neal Koblitz and Alfred Menezes, beginning with “Another Look at ‘Provable Security.’”¹ One of the main issues at stake here is the degree of assurance that provable security provides. Researchers have discovered many attacks on cryptographic schemes that were previously proven to be secure. However, we must consider that no science provides an absolute guarantee of the truth of its results, and that sciences evolve over time. Indeed, compared to more mature sciences, the field of cryptography, while developing rapidly, is still new. Moreover, the application of any scientific discipline always lags behind theoretical developments. As Einstein is reputed to have said, “In theory, theory and practice are the same. In practice they are not.” In this article, we illustrate some of the gaps that exist between cryptography’s theory and practice (see the “Cultures in Cryptography” sidebar).

Although researchers have applied provable security to almost all aspects of cryptography, we focus only on its application to analyzing the security of

symmetric encryption schemes used in Internet Protocol security (IPsec), Secure Sockets Layer (SSL)/Transport Layer Security (TLS), and Secure Shell (SSH) protocols. However, much of what we say here applies more widely to other aspects of cryptography and analysis techniques, such as formal methods and universal composability.

Two Milestones in Provable Security

Provable security dates back to 1949 when Claude Shannon used information-theoretic concepts to prove the *perfect secrecy* of the one-time pad in a paper titled “A Mathematical Theory of Cryptography.”² Perfect secrecy is an information theoretic notion of security that, informally, defines an encryption scheme to be perfectly secure if the ciphertext doesn’t leak information about the plaintext. Note that perfect secrecy makes no assumptions about the adversary’s computational capabilities, meaning that an exhaustive search on the encryption key would convey no additional information. The drawback of a perfectly secret encryption scheme is that its key space needs to be at least as large as its plaintext space. This creates significant issues for the practical deployment of such a scheme. Essentially, the problem of secure communication is replaced by the problem of generating, distributing, and securely destroying the large amounts of keying material necessary for the scheme. Indeed, such schemes haven’t been widely deployed despite their proven security properties.

In 1984, Shafi Goldwasser and Silvio Micali made



JEAN PAUL
DEGABRIELE
AND KENNETH
G. PATERSON
Royal
Holloway,
University
of London

GAVEN J.
WATSON
University
of Calgary

Cultures in Cryptography

It might seem strange that there can be such an obvious discrepancy between the theory and practice of cryptography. This gap might be in part due to the fact that different actors approach the field from completely different perspectives.

Theoreticians

First, consider a typical theoretical cryptographer's viewpoint. Judging by cryptographic research literature, one might surmise that many in the community view cryptography as a branch of theoretical computer science. With notable exceptions, the literature addresses theoretical questions as opposed to the real-world problems currently affecting widely used protocols. Although this focused approach has driven the field forward as a theoretical subject, it has done less to improve the security of currently used systems and makes the field seem less accessible to those in the applied community.

Practitioners

Provable security has become an important research area in modern cryptography but is still met with skepticism by many in the practical community. This is understandable considering the types of attack that we outline here. Practitioners might think provable security results provide an absolute statement of

security, especially if they're presented in such a manner. When they later discover that a scheme is insecure because of an attack outside the security model, this might damage their confidence in the whole enterprise of provable security.

In the practical community, there are at least two viewpoints to consider: the perspective of specification document writers and that of the implementers using these documents as a guide for writing code. Usually, writers must include some flexibility in a specification to allow for differences in implementations, allow interoperability, and account for the competing interests of the different parties contributing to the development process. Also, there's often a requirement to support certain options for backward compatibility. Unfortunately, this flexibility only increases the specifications' complexity and the consequent risk that an unfortunate feature interaction will lead to an attack.

Next, consider the practitioner implementing a system using cryptography. Because implementers base their implementations on the relevant specification documents and RFCs, they might not read theoretical research papers. With many options open to implementers in these documents, they might introduce a vulnerability. Without knowledge of the theory behind the choice being presented, implementers will have little idea of the effect a particular choice might have on security.

a breakthrough in provable security, introducing semantic security for public-key encryption.³ Semantic security adapts Shannon's perfect secrecy to the computational setting, considering only adversaries having bounded computational resources. Informally, semantic security states that no computationally bounded adversary given the ciphertext can predict anything about the plaintext any better than it can when it's not given the ciphertext. Goldwasser and Micali's original scheme was very inefficient in terms of computation and bandwidth, involving bit-by-bit encryption of plaintexts. However, we now have schemes that achieve semantic security efficiently, with security based on reasonable number-theoretic hardness assumptions. Thus, although weaker than perfect secrecy, semantic security doesn't pose the same practical limitations as perfect secrecy. Goldwasser and Micali's introduction of semantic security is widely recognized as the start of provable security as we know it today.

Understanding Security Proofs

Provable security, as its name implies, is about proving cryptographic schemes secure according to some notions of security. Such proofs commonly take the form of *reductions*, a concept borrowed from complexity theory. Reducing problem A to problem B means

that given access to an algorithm M_B that solves problem B , we can construct an algorithm M_A that employs M_B as a subroutine to efficiently solve problem A . Now suppose we can reduce solving problem P to breaking an encryption scheme E . Then, by contraposition, breaking E is at least as hard as solving P . Such a reduction would then represent a *conditional* proof of security for encryption scheme E on the basis that problem P is hard. However, we still need to clarify *break* and *hard*. We illustrate these concepts in the context of encryption schemes.

A break is normally defined through a game, in which one player (the challenger) gives well-defined capabilities to the other player (the adversary), and the adversary winning the game directly translates to a breach of a well-defined and well-motivated security notion. Semantic security for public-key encryption is an instance of such a notion but is cumbersome to use in proofs. Therefore, proofs often use the notion of *indistinguishability* (IND). Goldwasser and Micali showed that indistinguishability implies semantic security, and in "The Notion of Security for Probabilistic Cryptosystems," Micali and colleagues later proved the converse, showing that the two notions are equivalent (at least in the framework of polynomial-time reductions).⁴

The IND game for public-key encryption has two

phases. In the first phase, the challenger generates a public/private key pair for the scheme and gives the adversary the public key. The adversary performs its computation, selects two equal-length plaintexts m_0 , m_1 , and submits these to the challenger. The challenger randomly picks one of the two plaintexts, encrypts it, and gives the resulting challenge ciphertext to the adversary. In the second phase, the adversary performs further computation and then guesses which of its two chosen plaintexts was encrypted in the challenge ciphertext. The adversary wins the IND game if it guesses correctly. Because an adversary can always win this game with probability $1/2$ by guessing randomly, we consider the adversary successful (and the scheme broken) only if it wins the IND game with probability significantly greater than $1/2$. Notice that because the adversary has access to the public key, it could simply reencrypt the two plaintexts m_0 , m_1 in the second phase. If the encryption scheme was deterministic, then one of the two results would match the challenge ciphertext, and the adversary would win the IND game with probability 1. Thus, a semantically secure encryption scheme can't be deterministic.

Having described the adversary's goal in the IND game, we next classify its capabilities more carefully. Throughout our game, the adversary has the public key and can perform arbitrary encryptions. In the symmetric-key setting, we can't simply give the adversary the key, so we give the adversary access to an encryption oracle. The adversary submits arbitrary plaintexts to this oracle and in return receives ciphertexts encrypting those plaintexts. In public-key and symmetric-key settings, we then talk of a chosen plaintext attack (CPA). To capture a more powerful adversary, we also consider the chosen ciphertext attack (CCA) setting. Here, we give the adversary access to both the public key (or an encryption oracle for the symmetric-key setting) and a decryption oracle. The adversary sends ciphertext to the decryption oracle and in return obtains that ciphertext's decryption, or an indication that the decryption process failed. (However, we can't let the adversary query the challenge ciphertext to the decryption oracle, or winning the IND game would be trivial.) This oracle models the fact that, in practice, an adversary might have access to such a decryption capability, for example, through interaction with the scheme's users. It also reflects a general conservatism in provable security: to capture the widest set of possible attacks in our model, we try to make attackers as powerful as possible (without making it trivial for them to win). In addition to being given access to the oracles, the adversary can perform arbitrary computations. In this way, a security proof applies to all adversaries, not simply those

who behave in particular ways that might be limited by designers' imaginations.

However, we must limit the amount of resources consumed by the adversary. Otherwise, for example, the adversary could simply perform an exhaustive search for the private key, rendering our definitions vacuous. We use complexity theory and the framework of polynomial-time algorithms to assist us here. We introduce a security parameter k that can be thought of as defining the key size in the encryption scheme. Then an encryption scheme is said to be IND-CPA or IND-CCA secure if we can prove that no polynomial-time-bounded adversary can win the respective IND game with probability that is non-negligibly greater than $1/2$ in the parameter k . Here, a function $f(k)$ is said to be negligible if for every polynomial $p(k)$ there exists an N such that for all integers $k > N$, $f(k) < 1/p(k)$. Thus, a non-negligible function is one that grows faster than the inverse of some polynomial in k , and the scheme is secure if no polynomial-time adversary has an advantage that grows in this way.

The proof of such a statement takes the form of a reduction from a problem P to the problem of winning the particular scheme's IND-CPA or IND-CCA game with non-negligible probability. However, barring a major breakthrough in complexity theory, the existence of cryptographically useful problems P not having polynomial-time solutions remains in doubt. We can only prove our cryptographic scheme's security on the basis of the assumption that P really is sufficiently hard. Hence, the security proof is only conditional. For this reason, some consider the term *provable security* misleading or inappropriate and prefer the term *reductionist security*. Nonetheless, this approach is beneficial for several reasons. First, for many years, researchers have analyzed certain computational problems, such as integer factorization, the discrete logarithm problem over a finite field, and the problem of finding the shortest vector in a lattice, and these are now generally considered to be hard. More specifically, for certain choices of parameters, we can estimate the success probabilities and running times of algorithms for solving such problems. When a new cryptographic scheme is introduced, reducing a well-known and hard problem to its security transfers our confidence about the problem's hardness to the cryptographic scheme's security. Furthermore, because we can compare and contrast the hardness of problems through their complexity classes, we can use this as a criterion for comparing the security of cryptographic schemes.

Concrete security is a natural development of the above approach. Instead of working in the fairly abstract framework of polynomial-time algorithms and negligible functions, we try to directly relate any ad-

versary's success probability and running time to that of the algorithms to solve some assumed-to-be-hard problem P via the security reduction obtained in the proof. For example, suppose we can prove that for any adversary running in time t and having success probability $1/2 + \epsilon$ against the scheme, there's an algorithm to solve problem P that runs in time $f(t, \epsilon)$ and has success probability $g(t, \epsilon)$, for functions f and g . Then, by inverting this relationship, we can in principle use our knowledge about the current and projected state of the art in solving problem P along with our desired security level to select a concrete security parameter for the scheme. This approach also provides a means by which to judge the quality of proofs: a better proof is one that closely relates the resources an adversary consumed when breaking the scheme to those of an algorithm to solve the underlying problem P .

Another benefit of the provable security approach is that it enables the analysis of complex cryptographic schemes and protocols in terms of the security of simpler cryptographic primitives from which they're constructed. Such constructions might achieve more useful goals or meet stronger notions of security than their constituent components. By reducing a constructed scheme's security to that of its component primitives, we can achieve a higher security goal without introducing any further assumptions. Mihir Bellare and Chanathip Namprempre offered one such construction.⁵ They considered the sequential composition of a symmetric encryption scheme that is IND-CPA secure and a MAC (message authentication code) that is SUF-CMA (strongly unforgeable under chosen message attack) secure. Here, SUF-CMA is a standard security notion for a MAC. Bellare and Namprempre showed that such an encrypt-then-MAC construction yields an IND-CCA-secure encryption scheme. Hugo Krawczyk showed similar results for the analogous MAC-then-encrypt construction, for certain classes of encryption scheme.⁶

Some Limitations of Security Proofs

Provable security has obvious applications to cryptographic practice, such as protocol design and scheme selection. The existence of a correct security proof should be an important factor in selecting cryptographic schemes and adopting constructions for use in real-life protocols. However, we should take care when interpreting provable security results for practice. To reiterate, a security proof isn't an absolute guarantee of security. Rather, proofs are conditional, and security is guaranteed only as long as the underlying assumptions hold. But the reasons why caution is needed extend beyond this point.

In particular, security claims might not completely capture what should be considered a break in real

life. For instance, indistinguishability implies that the adversary can't distinguish between the encryptions of plaintexts of its choice, given that they're of equal length. However, in real life, an attacker might be able to predict something about the plaintext through the ciphertext's length unless we take extra precautions to prevent such traffic analysis. In some applications, this might constitute a confidentiality breach (see "Attacking and Repairing the WinZip Encryption Scheme" for a concrete example in the context of data compression⁷). In addition, the basic indistinguishability definition says nothing about denial of service or attacks based on message replays. Thus, we must be aware of what security is guaranteed by the proof as opposed to what our security requirements are.

Furthermore, the capabilities we gave the adversary in the security game might not accurately reflect those a real-world adversary possesses. We offer several examples of practical attacks in which adversaries can glean small amounts of information about the plaintext format from the manner in which the decryption process proceeds (or fails to proceed). In some cases, this information can be leveraged to enable real-world adversaries to extract plaintext from a challenge ciphertext. This information is rarely included in formal security models, but our examples show the importance of considering this kind of leakage. More generally, side-channel attacks use extra information produced by an implementation to mount an attack. These types of attacks might not be captured by traditional formal security analyses because they lie outside the scope of the security model. To obtain meaningful real-world security guarantees from a security proof, it's essential that the security model used accurately captures all the powers available to real-world adversaries.

Furthermore, designing and implementing secure cryptosystems—and combining various cryptographic primitives—are difficult tasks with many different aspects to consider. Vulnerabilities might be introduced inadvertently because of a fundamental system design flaw or by a programmer during implementation. Specification documents, if available, typically give implementers some flexibility, but without proper and detailed guidance, their choices can affect a system's security dramatically. Even a slight change to a provably secure scheme can render the system insecure. So how do we know if an implementation is an accurate translation of what has been modeled and proven secure?

Applying Provable Security to Secure Communications Protocols

Here, we highlight problems that arise when applying provable security results. Our examples focus on

symmetric cryptography as used in SSL/TLS, IPsec, and SSH. These examples make good case studies because the protocols are widely deployed and appear to be simple enough to be amenable to provable security analysis.

SSL/TLS

SSL is perhaps the most popular secure network protocol. One of its main uses is to secure credit-card data for Internet purchases. SSL was originally developed in the mid 1990s by Netscape and has since been adopted by the Internet Engineering Task Force (IETF) and renamed TLS. The current version of the protocol, TLS 1.2, is defined in RFC 5246. SSL/TLS consists of a number of subprotocols; here, we focus on the SSL/TLS Record Protocol, which provides confidentiality and integrity services for all SSL/TLS messages. At a high level, the Record Protocol uses a MAC-then-encrypt construction. This means that the message to be sent first has a MAC calculated on it and then encryption is performed over the concatenation of the message and the MAC tag. In addition, if a block cipher in CBC (cipher block chaining) mode is used, then the RFC specifies that we must add padding after the MAC is added and before the encryption. The padding must follow a particular format, and the RFC requires that this format be checked again after decryption.

Now consider what's known about encryption schemes following a MAC-then-encrypt structure from a provable security perspective. Bellare and Namprempe performed a formal analysis of generic compositions of a MAC and an IND-CPA-secure symmetric encryption scheme.⁵ They studied three constructions in detail: encrypt-and-MAC, encrypt-then-MAC, and MAC-then-encrypt. Of these, they showed that encrypt-then-MAC is the only one that can achieve IND-CCA security (starting with appropriate notions of security for the encryption and MAC components). Indeed, they showed that the MAC-then-encrypt construction isn't IND-CCA secure by exhibiting an SUF-CMA-secure MAC scheme and an IND-CPA secure encryption scheme whose MAC-then-encrypt composition can be broken in the IND-CCA attack model. Their analysis tells us nothing positive about the specific MAC-then-encrypt case used in SSL/TLS. Krawczyk further analyzed the MAC-then-encrypt composition, showing that if we use CBC mode of a good block cipher as the encryption component, then a MAC-then-encrypt construction will provide a secure channel.⁶ A secure channel is weaker than IND-CCA security but provides a positive result about the construction's security. In fact, Krawczyk's results can be improved to show that MAC-then-encrypt achieves IND-CCA

security when the encryption component is instantiated using CBC mode of a good block cipher.

So we now have a security proof for the MAC-then-encrypt construction using CBC mode. Is this enough to claim security of the SSL/TLS Record Protocol in practice? Even setting aside implementation questions, the answer, unfortunately, is no. Brice Canvel and colleagues' attack against the MAC-then-encrypt construction used in SSL/TLS worked in practice against OpenSSL, one of the most widely used implementations.⁸ Their attack exploits the way SSL/TLS handles padding during encryption. Recall that when SSL/TLS uses CBC mode, some padding must be added to the message before encryption. This padding's format should also be checked during decryption, with an error message returned over the Record Protocol if the format is incorrect. Similarly, an error message is returned if the MAC fails to verify. These errors are "fatal"—the SSL/TLS channel is destroyed if either of the errors arises. Because of the construction, it's natural to check the padding before the MAC when decrypting. Canvel and colleagues' attack exploits the small timing difference that arises between the appearance of a padding error and MAC error to mount a special type of side-channel attack called a *padding oracle attack*. This type of attack can recover plaintext from arbitrary ciphertext blocks. Moreover, if a fixed plaintext is repeated across many sessions (for example, a password), then the attack can be iterated to boost its success probability even though the channel is destroyed at each attempt.

Canvel and colleagues' attack seems to contradict Krawczyk's security proof. How does the contradiction arise? Krawczyk's proof is mathematically correct, but his model doesn't accurately capture the way SSL/TLS works in practice. Indeed, the particular construction that Krawczyk studied doesn't include padding: it simply assumes that the plaintext lengths are already suitable for applying a block cipher in CBC mode. Also, his analysis doesn't consider the possibility of distinguishable error outputs. This means that although Krawczyk's analysis is correct and rules out many possible attacks against SSL/TLS, it can't be applied directly to the protocol as specified. Instead, we need a more accurate model that reflects the fact that SSL/TLS really uses a MAC-then-pad-then-encrypt construction. Kenneth Paterson and Gaven Watson have taken initial steps in this direction in "Immunising CBC Mode against Padding Oracle Attacks: A Formal Security Treatment," in which they study the CBC mode's IND-CPA security in a situation in which the attacker has access to padding error information.⁹ However, extending this type of analysis to the IND-CCA case for the specific construction used in SSL/TLS remains an open problem. We still don't

have strong, formally proven guarantees for the SSL/TLS Record Protocol.

Note that the IETF community's reaction to Canvel and colleagues' SSL/TLS attacks wasn't to change the protocol to use a construction known to be IND-CCA secure in the face of padding oracle attacks (for example, by using a different padding scheme or by using a different construction entirely). Instead, RFC 5246 recommends that implementers ensure the MAC and padding errors aren't distinguishable by content or timing. This prevents the attack and requires minimal code changes. Thus, a pragmatic but unproven fix has been the solution of choice, reflecting the constraints and compatibility issues the IETF community faces when changing such a significant protocol.

IPsec

IPsec is a protocol suite offering security services at the TCP/IP stack's IP layer. IPsec is most widely deployed to build virtual private networks (VPNs) and secure remote access solutions. Its major components are the AH (authentication header), ESP (encapsulating security payload), and IKE (Internet key exchange) protocols specified in RFCs 4302, 4303, and 4306, respectively. AH provides data-origin authentication and replay protection services using a MAC in combination with sequence numbers and a replay window. ESP originally provided a data confidentiality service but now provides data-origin authentication as well. IKE provides an automated security capability negotiation and key-management service. Because of IPsec's flexibility, users have various options for how to combine these protocols to build a secure channel that suits their needs.

Jean Paul Degabriele and Kenneth Paterson recently demonstrated a series of plaintext-recovery attacks against various IPsec configurations in which AH is first applied to an IP datagram and the result is then protected by ESP in encryption-only mode.¹⁰ This is yet another instantiation of the MAC-then-encrypt construction, analyzed by Bellare and Namprempre⁵ and Krawczyk.⁶ These attacks work regardless of the block cipher or MAC algorithm, as long as we use the block cipher in CBC mode. As with Canvel and colleagues' SSL/TLS attack, Degabriele and Paterson's attacks exploit several features from the IPsec realization of the basic MAC-then-encrypt composition, which aren't captured in Krawczyk's security model. These include fields in the AH-protected packet that aren't covered by the MAC algorithm, extra unauthenticated padding fields in the plaintext before encryption (namely, traffic flow confidentiality [TFC] padding bytes, encryption padding bytes, and the next-header byte), and the fact that IPsec should check certain padding fields for correctness after de-

ryption. Here, incorrect formatting leads to a packet drop, whereas correct formatting leads to further processing, which can be arranged to result in a response message being sent on the secure channel. The presence (or absence) of such a message can reveal a small amount of information about the plaintext format, which can in turn be leveraged to recover arbitrary amounts of plaintext in a reliable manner. One of Degabriele and Paterson's attacks also exploits the manner in which IPsec's processing interacts with IP fragment processing to generate error messages that the attacker can detect.

These attacks on IPsec—as with those on SSL/TLS—demonstrate that an attacker can exploit error information not normally considered in formal analyses to mount attacks against provably secure cryptosystems. Formal analyses that don't consider such error information will typically fail to capture all the subtleties in the implementation of a secure channel protocol, and this can lead to attacks.

Another important distinction between Krawczyk's theoretical model and the way in which cryptography operates in IPsec relates to atomicity. Krawczyk's model treats the MAC-then-encrypt construction as an atomic operation (although a counterexample given in his article highlights that MAC-then-encrypt might not be secure when an intermediate encoding step is introduced). However, in IPsec, the MAC-then-encrypt configurations are realized by combining two separate protocols, AH and ESP, with each protocol performing its own processing, and with the possibility of additional processing between these steps.

SSH

Finnish researcher Tatu Ylönen originally designed SSH in the mid 1990s to replace insecure remote logins after his university's network was targeted by some password-sniffing attacks. In 2006, a collection of RFCs specified SSH version 2. We're interested in RFC 4253, which specifies the SSH Binary Packet Protocol (BPP)—the part of SSH responsible for ensuring messages' confidentiality and integrity. The BPP follows an encode-then-encrypt-and-MAC construction. The payload message is first encoded and then encrypted; the MAC is calculated on the encoded message together with a sequence number. A ciphertext packet consists of the concatenation of the encrypted message and the MAC tag. The encoding scheme adds three fields to the payload message: a packet-length field, a padding-length field, and some padding bytes. The packet-length field specifies the total length of what follows, that is, the combined length of the padding-length field, the payload message, and the padding field. It's ostensibly encrypted

to protect against traffic analysis. This has a significant effect on the protocol's security.

Mihir Bellare and colleagues performed a formal analysis of the SSH BPP.¹¹ Because of a distinguishing attack by Wei Dai that exploits the use of initial packet chaining when using CBC mode encryption,¹² Bellare and colleagues weren't able to directly prove security for the SSH BPP as defined in RFC 4253. Instead, they proposed several minor SSH BPP variants and proved them secure in an extended version of the IND-CCA model, with the new model taking into account the stateful nature of SSH decryption.

As with our previous examples, the SSH BPP turned out to be vulnerable to attack, despite having a mathematically correct security proof. Martin Albrecht and colleagues discovered plaintext-recovery attacks that exploit the SSH BPP's use of an encrypted packet-length field, its reliance on CBC mode, and the attacker's ability to send ciphertext data in small chunks and observe how the recipient reacts.¹³ The attack principles are simple. An attacker observes a ciphertext and chooses one block to attack. The attacker sends this target block to the recipient in such a way that the recipient interprets it to be the start of a new packet. The recipient must immediately decrypt this block to retrieve the packet-length field, to know how much data it must wait for before it receives and verifies the MAC. The attacker then proceeds by sending random blocks one at a time until the recipient outputs a MAC error. By counting how many random blocks have been sent, the attacker can deduce the new packet's packet-length field and, by the properties of CBC mode, deduce the corresponding bits in the target plaintext block. In practice, this attack is complicated by checks performed on the packet-length field once the recipient covers it. Albrecht and colleagues implemented variants of this attack against OpenSSH, one of which recovers 32 bits of plaintext with probability 2^{-18} .

This attack can be applied to one of the provably secure variants of the SSH BPP proposed by Bellare and colleagues. So what went wrong? First, their analysis assumes that ciphertexts are self-describing in terms of their lengths. In reality, we see that recipients must decrypt the first block of a packet as soon as they receive it to obtain the packet length. RFC 4253 actually states that

[i]mplementations SHOULD decrypt the length after receiving the first 8 (or cipher block size, whichever is larger) bytes of a packet.

In addition, Bellare and colleagues' model doesn't allow for the possibility that the amount of data needed to complete the decryption process is governed by data produced *during* the decryption process. Second,

we note that in their analysis, ciphertexts and plaintexts are handled as atomic strings. In contrast, Albrecht and colleagues' attacks exploit the fact that an attacker can send data in small chunks to the recipient. Many implementations use a buffer to store data until it's needed, but Bellare and colleagues' analysis doesn't model this.

Kenneth Paterson and Gaven Watson recently provided a new formal analysis of the SSH BPP using counter-mode encryption, with the explicit intention of addressing the shortcomings of the previous analysis.¹⁴ Paterson and Watson begin by defining a new version of SSH-CTR that accurately captures how the SSH BPP with counter-mode encryption is defined in the RFCs and coded in practice in OpenSSH and other implementations. They also extended the previous security model to account for the manner in which the SSH BPP buffers as-yet-unprocessed ciphertext bytes, and to let the attacker deliver ciphertext to a decryption oracle in a byte-by-byte fashion. They then proved that their new definition of SSH-CTR is secure in the new model. Because this new analysis more accurately captures how SSH is defined in the RFCs and how it's implemented, we can be more confident that this provable security result is more meaningful in practice.

The Future

Provable security is maturing into a very useful tool that can and should play an important role in the design and analysis of systems using cryptography. However, we've seen that protocols can be vulnerable to attacks outside the scope of the existing formal analyses. Nonetheless, such analyses can rule out large classes of attack. This leaves us with two major questions: How can we make the formal analyses more closely related to real implementations? And how can we better integrate formal analyses with the design process?

Making Theory More Applicable

In the SSH example, we saw that it's possible to perform analyses that accurately capture how schemes

Provable security is maturing into a very useful tool that can and should play an important role in the design and analysis of systems using cryptography.

operate in practice. Researchers are actively pursuing further work in this area, including the examination of how hardware and software side-channel attacks can be captured by formal analyses. For example,

leakage-resilient cryptography is a particularly active research area that attempts to reason about security when the attacker has access to some secret information gained using a side channel.

When considering how to make the theory more practical, we still have at least one fundamental question to answer: How do we know exactly which protocol features are critical to its security and therefore must be included in the formal analysis? Many fine-grained details in cryptographic implementations might ultimately play a significant role in the protocol's overall security. As we've shown, omitting one of these details from a formal analysis can have a large effect on how that analysis applies in practice. On the other hand, Paterson and Watson's analysis of SSH attempts to model exactly how the RFCs specify the SSH BPP and how it's implemented in OpenSSH. Because of the model's detail and accuracy, we now have more confidence that no significant attack vectors were omitted. That said, Paterson and Watson's approach involves manual code inspection and proof generation. In the longer term, we hope that developers will create automated tools to assist in this kind of task. However, even this might not be sufficient—their analysis doesn't consider any issues arising from compression, which can play an important role in determining protocol security.⁷

Phillip Rogaway and Till Stegers have taken an alternative to Paterson and Watson's approach.¹⁵ They consider a protocol to consist of two parts: a partially specified protocol (PSP) and the protocol details (PD). The PSP will be the protocol's cryptographic core and is strictly defined by the protocol's specifications documents. The PD encompasses additional features necessary for an implementation but left open for the implementer to choose. Again, these additional choices sometimes introduce security vulnerabilities. To address this, Rogaway and Stegers give the adversary control over the PD, then attempt to prove the PSP's security. The idea is that if we can prove the PSP is secure when the adversary has control over the PD, then the full protocol (PSP, PD) must be secure in practice for any PD implementation. This is an intriguing approach that we expect to see further developed.

Bringing Theory into Practice

Today, formal protocol analyses are predominantly performed in a reactive manner: typically, someone designs and specifies a protocol, and only later does someone perform a formal analysis. Along with this approach, we've seen the rise of a model-attack-remodel cycle, with the models gradually being upgraded to more closely reflect practice in light of new attacks. It's reminiscent of cryptography's break-fix cycle but has the advantage that, at each iteration, the

security proof rules out a larger class of attacks. In an ideal world, we'd perform proactive formal analyses as part of the design process, before systems become widely deployed and therefore harder to change. By bringing together analysis and design, we can get a better idea of how various options might affect system security. Specification documents should include formal analysis results, giving the implementer a clear rationale for choices made in the design phase. However, this requires careful presentation, with provable security results presented in a simple and understandable manner for the consumption of practitioners.

In our view, better communication is key to uniting the theoretical and applied communities in cryptography. Without this, the two communities will continue to pull in opposite directions. By engaging both theoreticians and practitioners in a common endeavor—that of improving the security of real systems—we will foster a greater understanding of what provable security results actually mean and speed their adoption by practitioners. We hope that this article provides a small step in this direction. □

Acknowledgments

Vodafone Group Services Limited, a Thomas Holloway Research Studentship, and the Strategic Educational Pathways Scholarship Scheme (Malta), partly financed by the European Union–European Social Fund support Degabriele's research. An EPSRC Leadership Fellowship, EP/H005455/1, supports Paterson's research. An EPSRC Industrial CASE studentship sponsored by BT Research Laboratories supports Watson's research.

References

1. N. Kobitz and A. Menezes, "Another Look at 'Provable Security,'" *J. Cryptology*, vol. 20, no. 1, 2007, pp. 3–37.
2. C. Shannon, "Communication Theory of Secrecy Systems," *Bell System Technical J.*, vol. 28, no. 4, 1949, pp. 656–715.
3. S. Goldwasser and S. Micali, "Probabilistic Encryption," *J. Computer Systems Science*, vol. 28, no. 2, 1984, pp. 270–299.
4. S. Micali, C. Rackoff, and B. Sloan, "The Notion of Security for Probabilistic Cryptosystems," *CRYPTO 1986*, LNCS 263, Springer, 1986, pp. 381–392.
5. M. Bellare and C. Nampreppe, "Authenticated Encryption: Relations among Notions and Analysis of the Generic Composition Paradigm," *ASIACRYPT 2000*, LNCS 1976, Springer, 2000, pp. 531–545.
6. H. Krawczyk, "The Order of Encryption and Authentication for Protecting Communications (or How Secure Is SSL?), " *CRYPTO 2001*, LNCS 2139, Springer, 2001, pp. 310–331.

7. T. Kohno, "Attacking and Repairing the WinZip Encryption Scheme," *ACM Conf. Computer and Comm. Security*, ACM Press, 2004, pp. 72–81.
8. B. Canvel et al., "Password Interception in a SSL/TLS Channel," *CRYPTO 2003*, LNCS 2729, Springer, 2003, pp. 583–599.
9. K.G. Paterson and G.J. Watson, "Immunising CBC Mode against Padding Oracle Attacks: A Formal Security Treatment," *SCN*, LNCS 5229, Springer, 2008, pp. 340–357.
10. J.P. Degabriele and K.G. Paterson, "On the (In)security of IPsec in MAC-then-Encrypt Configurations," *Proc. 17th ACM Conf. Computer and Comm. Security (CCS 10)*, ACM Press, 2010, pp. 493–504.
11. M. Bellare, T. Kohno, and C. Namprempre, "Breaking and Provably Repairing the SSH Authenticated Encryption Scheme: A Case Study of the Encode-then-Encrypt-and-MAC Paradigm," *ACM Trans. Information and Systems Security*, vol. 7, no. 2, 2004, pp. 206–241.
12. W. Dai, "An Attack Against SSH2 Protocol," 6 Feb. 2002; www.ietf.org/mail-archive/text/secsh/2002-02. mail.
13. M.R. Albrecht, K.G. Paterson, and G.J. Watson, "Plaintext Recovery Attacks against SSH," *IEEE Symp. Security and Privacy*, IEEE CS Press, 2009, pp. 16–26.
14. K.G. Paterson and G.J. Watson, "Plaintext-Dependent Decryption: A Formal Security Treatment of SSH-CTR," *EUROCRYPT 2010*, LNCS 6110, Springer, 2010, pp. 345–361.
15. P. Rogaway and T. Stegers, "Authentication without Elision: Partially Specified Protocols, Associated Data, and Cryptographic Models Described by Code," *22nd Computer Security Foundations Symp. (CSF 09)*, IEEE CS Press, 2009, pp. 26–39.

Jean Paul Degabriele is a PhD student at Royal Holloway, University of London. His research interests include cryptography and network security. Degabriele has an MSc in information security from the University of London. Contact him at j.p.degabriele@rhul.ac.uk.

Kenneth G. Paterson is a professor of information security at Royal Holloway, University of London. His research interests include cryptography and information security. Paterson has a PhD in mathematics from the University of London. He's a fellow of the IMA and a Journal of Cryptology editorial board member. Contact him at kenny.paterson@rhul.ac.uk.

Gaven J. Watson is a postdoc at the University of Calgary. His research interests are in cryptography and network security. Watson has a PhD in mathematics from Royal Holloway, University of London. Contact him at gavenjwatson@gmail.com.



Selected CS articles and columns are also available for free at <http://ComputingNow.computer.org>.



handles the details
so you don't have to!

- Professional management and production of your publication
- Inclusion into the IEEE Xplore and CSDL Digital Libraries
- Access to CPS Online: Our Online Collaborative Publishing System
- Choose the product media type that works for your conference:
Books, CDs/DVDs, USB Flash Drives, SD Cards, and Web-only delivery!

Contact CPS for a Quote Today!

www.computer.org/cps or cps@computer.org



IEEE  computer society